


# Clustering



INF 553:  
Foundations and Applications of Data Mining

# Roadmap

- ◆ Problem, types, and distance functions
- ◆ Algorithms:
  - ◆ Hierarchical clustering
  - ◆ Point assignment
    - ◆ K-means
    - ◆ BFR: extend k-means to handle large data set
    - ◆ CURE

# Learning Approaches

## Supervised Learning

- ◆ The **training data** is **annotated** with information to help the learning system
- ◆ Eg. the class for each instance

## Unsupervised Learning

- ◆ The training data is **not annotated** with any extra information to help the learning system
- ◆ Eg. clustering of data

## Semi-Supervised Learning

# High Dimensional Data

## High dim. data

Locality sensitive hashing

Clustering

Dimensionality reduction

## Graph data

PageRank, SimRank

Network Analysis

Spam Detection

## Infinite data

Filtering data streams

Web advertising

Queries on streams

## Machine learning

SVM

Decision Trees

Perceptron, kNN

## Apps

Recommender systems

Association Rules

Duplicate document detection

# High Dimensional Data


- ◆ Given a cloud of data points we want to understand its **structure**
- ◆ Group points into “**clusters**” according to some **distance measure**.



# The Problem of Clustering


- ◆ Given a **set of points** that belong to some **space**, with a notion of **distance** between points
- ◆ **Group the points** into some number of *clusters*, so that:
  - Members of **a cluster** are **close/similar** to each other
  - Members of **different clusters** are **dissimilar**
- ◆ **Usually:**
  - Points are in a high-dimensional space
  - Similarity is defined using a **distance measure**
    - Euclidean, Cosine, Jaccard, edit distance, ...

# Example: Clusters & Outliers





# Clustering is a hard problem!





# Why is it hard?

- ◆ Clustering in **two dimensions** looks easy
- ◆ Clustering **small amounts** of data looks easy
  - And in most cases, looks are **not** deceiving
- ◆ BUT: Many applications involve not 2, but 10 or 10,000 dimensions
- ◆ **High-dimensional spaces look different:** Almost all pairs of points are at about the same distance
  - **Curse of dimensionality:** In high dimensions, almost all pairs of points are **equally far away** from one another; almost any **two vectors are orthogonal**.

# Clustering Sky Objects: SkyCat

- ◆ A catalog of 2 billion “sky objects” represents objects by their radiation in 7 dimensions (frequency bands)
- ◆ **Problem:** Cluster into similar objects, e.g., galaxies, nearby stars, quasars, etc.
- ◆ Sloan Digital Sky Survey is a newer, better version of this



# Clustering Problem: Songs

- ◆ **Intuitively:** Music divides into categories, and customers prefer a few categories
  - But what are categories really?
- ◆ Represent a song by a **set of customers** who bought it
- ◆ **Similar songs** have similar sets of customers, and vice-versa.

# Clustering Problem: Songs

## Space of all songs:

- ◆ Think of a space with one dimension for each customer
  - Values in a dimension may be 0 or 1 only
  - A **song is a point** in this space  $(x_1, x_2, \dots, x_k)$ , where  $x_i = 1$  iff the  $i^{\text{th}}$  customer bought the song
- ◆ For Amazon, the dimension is tens of millions
- ◆ **Task:** Find clusters of similar songs.

# Clustering Problem: Documents

## Finding topics:

- ◆ Represent a document by a **vector**  $(x_1, x_2, \dots, x_k)$ , where  $x_i = 1$  iff the  $i^{\text{th}}$  word (in some order) appears in the document
  - It actually doesn't matter if  $k$  is infinite; i.e., we don't limit the set of words
- ◆ Representing documents by sets of shingles in another example
- ◆ **Documents with similar sets of words or same shingles may be about the same topic.**


# Jaccard, Euclidean and Cosine Distance

- ◆ Different ways of representing documents (**as sets of words or shingles**) lead to different distance measures
- ◆ Document = **set** of words
  - **Jaccard distance**
- ◆ Document = **point** in space of words
  - $(x_1, x_2, \dots, x_n)$ , where  $x_i=1$  iff word  $i$  appears in doc
  - **Euclidean distance**
- ◆ Document = **vector** in space of words
  - Vector from origin to  $(x_1, x_2, \dots, x_n)$
  - **Cosine distance.**

# Euclidean Distance

- ◆ Measures distance of two points in Euclidean space

$$d([x_1, x_2, \dots, x_n], [y_1, y_2, \dots, y_n]) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$






# Cosine Distance

- ◆ Similarity = Cosine of angle btw vectors: A & B
  - Numerator is the **dot product** of vectors A and B
  - Denominator is the product of the **Euclidean distance** of each vector from the origin (length of the vector)

- ◆ distance = 1- Cosine(A, B)



$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

## Example-Cosine Distance


- ◆ Think of a point as a vector from the origin (0,0,...,0) to its location
- ◆ Two points' vectors make an angle, whose cosine is the normalized dot product of the vectors:  
 $p1.p2 / |p2| |p1|$
- ◆ **Example:**  $p1=00111$ ;  $p2=10011$
- ◆  $p1.p2=2$
- ◆  $|p2|=|p1|=\sqrt{3}$
- ◆  $\cos(\theta)=2/3$ ;  $\theta$  is about 48 degrees.

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

## For Sets, Jaccard Similarity

D1: The quick brown fox jumped over the dog.

D2: The brown fox jumps and the dog is there too.



$$\text{Jaccard} = |S \cap T| / |S \cup T| = 6/12 = 0.5$$

## Jaccard *Distance*

$$1 - (\text{Jaccard Similarity})$$


# Hamming Distance

- ◆ For two bit vectors, distance btw x and y =
  - # of corresponding bits that differ
- ◆  $x = \mathbf{10101}$ ,  $y = \mathbf{11110}$ 
  - $\text{Hamming}(x, y) = 3$

# Edit Distance

- ◆ Use when **points are strings**
- ◆ Distance between strings  $x = x_1x_2\dots x_n$  and  $y = y_1y_2\dots y_m$
- ◆ **Smallest number of insertions and deletions of single characters that will convert  $x$  to  $y$**
- ◆ Example:
  - $x = abcde$  and  $y = acfdeg$
  - To convert  $x$  to  $y$ : Delete  $b$ , insert  $f$  after  $c$ , insert  $g$  after  $e$
  - Edit distance = 3

# Roadmap

- ◆ Problem, types and distance functions
- ◆ Hierarchical clustering 
- ◆ Point assignment
  - ◆ K-means
  - ◆ BFR
  - ◆ CURE



# Overview: Methods of Clustering


## ◆ Hierarchical:

### ➤ Agglomerative (bottom up):

- Initially, each point is a cluster
- Repeatedly combine the two “nearest” clusters into one


### ➤ Divisive (top down):

- Start with one cluster and recursively split it



## ◆ Point assignment:

- Maintain a set of clusters
- Points belong to “nearest” cluster.




## Clustering Strategies (cont.)

Also distinguish clustering algorithms by:

- ◆ Whether the algorithm assumes a **Euclidean space** or uses some other **distance measure**
- ◆ Whether the algorithm assumes data are **small enough** to fit into memory.

# Hierarchical Clustering (Agglomerative)

- ◆ **Key operation:**  
**Repeatedly combine**  
**two nearest clusters**



- ◆ **Three important questions:**
  - **1)** How do you represent a **cluster of more than one** point?
  - **2)** How do you determine the “**nearness**” of clusters?
  - **3)** When to **stop combining** clusters?

# Hierarchical Clustering

- ◆ **Key operation:** Repeatedly combine two nearest clusters
- ◆ **(1) How to represent a cluster of many points?**
  - **Key problem:** As you merge clusters, how do you represent the “location” of each cluster, to tell which pair of clusters is closest?
- ◆ **Euclidean case:** each cluster has a ***centroid*** = average of its (data) points
- ◆ **(2) How to determine “nearness” of clusters?**
  - Measure cluster distances by distances of centroids.

# Hierarchical Clustering

- ◆ Initially, a point is in a cluster by itself

How to pick and combine efficiently?

When to stop?


```
WHILE it is not time to stop DO  
  [ pick the best two clusters to merge;  
    combine those two clusters into one cluster;  
  ]  
END;
```

How to measure cluster distance?<sup>27</sup>

# Hierarchical Agglomerative (Bottom-Up) Clustering Algorithm

- ◆ **First assume the space is Euclidean**
- ◆ Represent cluster by its **centroid** or average of points in the cluster
- ◆ **Merging rule:**
  - the distance between two clusters is distance between their centroids
  - $\text{dist}(C1, C2)$  = distance of their centroids
    - Coordinate of centroid = avg of that of all points in the cluster
  - $C1: \{(1, 2), (2, 2)\}$ 
    - Centroid =  $(1.5, 2)$
- ◆ **merge two clusters at shortest distance.**


# Example: Hierarchical clustering



**Data:**

o ... data point

x ... centroid



**Dendrogram**<sup>29</sup>



# When to Stop Clustering Process

◆ Several approaches:

**1. May know how many clusters** there are in the data

➤ Have been told or some **intuitive number of clusters**

**2. Stop combining when best combination of existing clusters produces a cluster that is inadequate**


➤ E.g., Average distance between centroid and its points should be below some limit.

# Rules for Controlling Hierarchical Clustering: Picking Clusters to Merge

1. Find pair with **smallest distance between centroids** (previous)
2. Take distance between two clusters as **minimum of distances between any two points, one chosen from each cluster**
  - **Merge two clusters with minimum distance**
  - May result in entirely different clustering from distance-of-centroids
3. Take distance between two clusters to be **average distance of all pairs of points, one from each cluster**
  - **Merge two clusters with smallest average distance**
4. **Radius of cluster** = maximum distance between all points and the centroid
  - **Combine two clusters whose resulting cluster has lowest radius**
5. **Diameter of cluster** = maximum distance between any two points of the cluster
  - **Merge the clusters whose resulting cluster has smallest diameter.**

# Dendrogram

- ◆ Can obtain  $k$  clusters from result for desired  $k$ 
  - $k$  can be any value between 1 and  $n$




# And in the Non-Euclidean Case?

## What about the Non-Euclidean case?

- ◆ The only “locations” we can talk about are the points themselves
  - i.e., there is no “average” of two points
  
- ◆ **Approach 1:**
  - (1) How to represent a cluster of many points?
    - *clustroid* = (data)point “closest” to other points
  - (2) How do you determine the “nearness” of clusters?
    - Treat clustroid as if it were centroid, when computing inter-cluster distances.

# Clustroid



**Centroid** is the avg. of all (data)points in the cluster. This means centroid is an “**artificial**” point

**Clustroid** is an **existing** (data)point that is “closest” to all other points in the cluster.

# “Closest” Point?


## ◆ (1) How to represent a cluster of many points?

**clustroid** = point “closest” to other points

## ◆ Possible meanings of “closest”:

- Smallest maximum distance to other points
- Smallest average distance to other points in the cluster
- Smallest sum of squares of distances to other points

- For distance metric  $d$  clustroid  $c$  of cluster  $C$  is:  $\min_c \sum_{x \in C} d(x, c)^2$



# Defining “Nearness” of Clusters

## ◆ (2) How do you determine the “nearness” of clusters?

### ➤ Approach 1:

**Intercluster distance** = minimum of the distances between any two points, one from each cluster

### ➤ Approach 2:

Pick a notion of “**cohesion**” of clusters, *e.g.*, maximum distance from the **clustroid**

- Merge clusters whose *union* is most cohesive.



# Termination condition

## ◆ (3) When to stop merging

- **Approach 1:** Pick a number **k** upfront, and stop when we have **k** clusters
  - Makes sense when we know that the data naturally falls into **k** classes
- **Approach 2:** Stop when the next merge would create a cluster with low “**cohesion**”
  - i.e, a “bad” cluster.

# Cohesion

- ◆ Merge clusters whose *union* is most cohesive
- ◆ Approach 3.1: **Diameter** of the merged cluster = maximum distance between points in the cluster
- ◆ Approach 3.2: **Radius** = maximum distance of a point from **centroid** (or clustroid)
- ◆ Approach 3.3: Use a **density-based approach**
  - Density = number of points per unit volume
  - E.g., divide number of points in cluster by **diameter or radius of the cluster**
  - Perhaps use a power of the radius (e.g., square or cube).

# Example

◆ Consider a cluster of 4 points:

➤ abcd, aecdb, abecb, ecdab

◆ Their edit distances:

	aecdb	abecb	ecdab
abcd	3	3	5
aecdb		2	2
abecb			4

# Determine Clusteroid

- ◆ aecdb will be chosen as clusteroid
  - Located in “center” judged by all 3 measures

	aecdb	abecb	ecdab
abcd	3	3	5
aecdb		2	2
abecb			4

Point	Sum	Sum-sq	Max
abcd	11	43	5
aecdb	<b>7</b>	<b>17</b>	<b>3</b>
abecb	9	29	4
ecdab	11	45	5

# Complexity of Hierarchical Clustering

- ◆  $n$  data points
- ◆ At most  $n - 1$  step of merging
- ◆ Naive implementation, e.g., storing pairwise cluster distances in a matrix

	<b>C1</b>	<b>C2</b>	<b>C3</b>	<b>C4</b>
<b>C1</b>	0	2	3	2
<b>C2</b>		0	4	5
<b>C3</b>			0	3
<b>C4</b>				0

# Implementation

- ◆ Naïve implementation of hierarchical clustering:
  - At each step, compute pairwise distance between all pairs of clusters, then merge.
- ◆ Initially,  $O(n^2)$  for creating matrix and finding pair with minimum distance
- ◆ Subsequent merge,  
=> Overall complexity:  $O(n^3)$

# Improved Version

- ◆ Use priority queue (e.g., heap-based) instead of matrix
- 1. Compute pairwise dist. of all points:  $O(n^2)$
- 2. Build priority queue (time linear to size of queue), so:  $O(n^2)$
- 3. Each merge:
  - a) Remove entries for old clusters:  $2n * O(\log(n))$
  - b) Add entries for new cluster:  $n * O(\log(n))$

=> Overall complexity:  $O(n^2 \log(n))$

# Implementation

## ◆ Naïve implementation of hierarchical clustering:

➤ At each step, compute pairwise distances between all pairs of clusters, then merge

- Initially,  $O(n^2)$  for creating matrix and finding pair with minimum distance
- Subsequent merge,


=> Overall complexity:  $O(n^3)$

◆ Careful implementation using **priority queue** can reduce time to  $O(N^2 \log N)$


➤ **Still too expensive for really big datasets that do not fit in memory.**



# Roadmap


- ◆ Problem, types and distance functions
- ◆ Hierarchical clustering
- ◆ Point assignment 
  - K-means
  - BFR
  - CURE

# K-Means Clustering Algorithm




- ◆ User specifies a target number of clusters ( $k$ )
- ◆ Place randomly  $k$  **cluster centers**
- ◆ For each datapoint, attach it to the nearest cluster center
- ◆ For each center, find the **centroid** of all the datapoints attached to it
- ◆ Turn the centroids into cluster centers
- ◆ Repeat until the sum of all the datapoint distances to the cluster centers is minimized


# K-Means Clustering (1)




## K-Means Clustering (2)




# K-Means Clustering (3)




# K-Means Clustering (4)



# K-Means Clustering (5)




# K-Means Clustering (6)





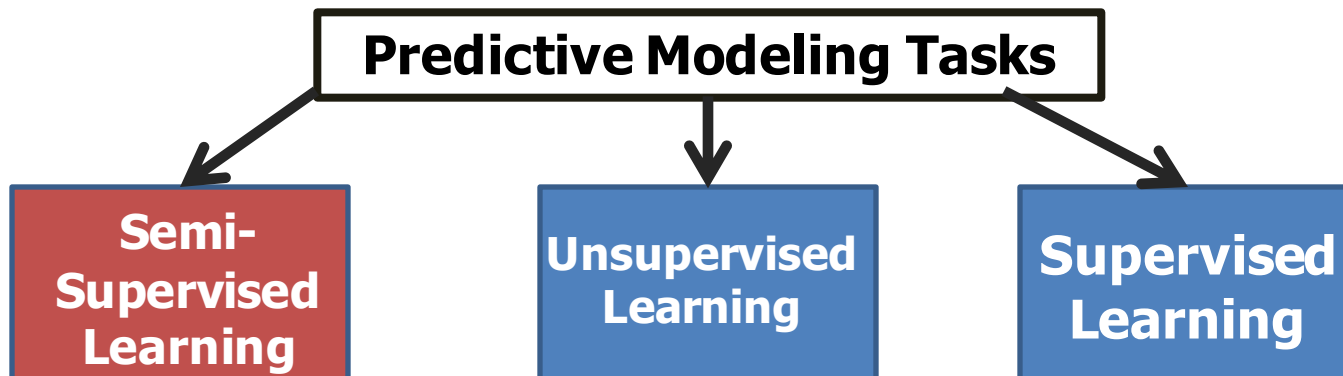
# Clustering Methods

- ◆ Hierarchical clustering
  - Attach datapoints to root points
- ◆ K-Means clustering
  - Centroid-based
- ◆ Density-based methods
  - Clusters contain a minimal number of datapoints



# Summary of Concepts in Clustering

## Clustering vs. Classification



### ◆ Classification is **supervised**

- class labels are provided;
- learn a **classifier to predict class labels** of novel/unseen data

### ◆ Clustering is **unsupervised** or **semi-supervised**;

- No class label is give
- Understand the structure underlying your data.

# Summary of Concepts in Clustering

## ◆ Clustering

## ◆ Algorithms:

### ◆ Hierarchical clustering

- centroid
- clustroid
- Dendrogram

### ◆ Point assignment

- K-means: cluster centers, centroids
- BFR: extend k-means to handle large data set
- CURE.