

# INF 553: Foundations and Applications of Data Mining

Map-Reduce

**Dr. Anna Farzindar** farzinda@usc.edu

# Map-Reduce: A diagram

Input

### MAP:

Read input and produces a set of key-value pairs

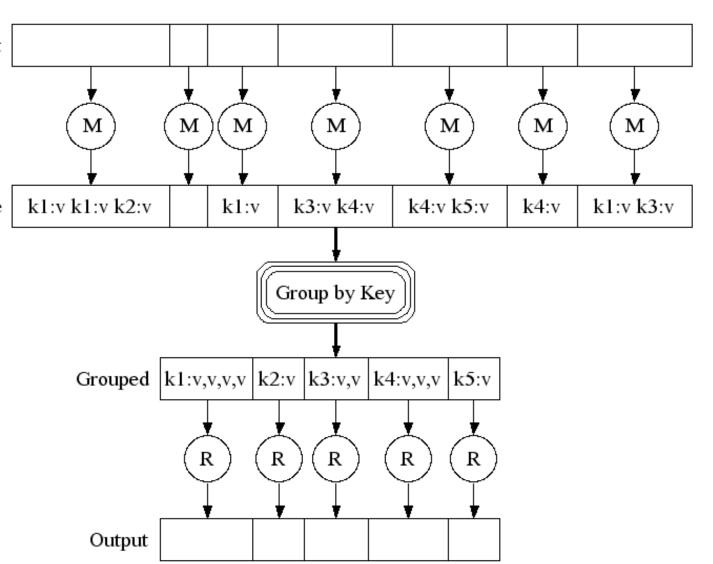
Intermediate

### Group by key:

Collect all pairs with same key
(Hash merge, Shuffle, Sort, Partition)

### Reduce:

Collect all values belonging to the key and output



# Example (Exercise 2.3.1 in book)

Design MapReduce algorithms to take a very large file of integers and produce as output:

**♦** (a) The largest integer

Map task: ?

Reduce tasks?

# Notes on example (a) (a) The largest integer

- **◆** Each Map task processes a chunk of the input file
- **◆ Map task produces (integer, 1) of the largest value in that chunk as key, value pair**
- Grouping by key identifies duplicates
- **◆** How many Reduce tasks?
- **◆** Single reduce task: produces (integer, 1) of largest value
- ◆ What about multiple reduce tasks?
  - In the shuffle step, keys could be **sorted by range**, so only look at **output** from the Reduce stage that has the **highest range of keys**.

# Exercise 2.3.1 (cont.)

Design MapReduce algorithms to take a very large file of integers and produce as output:

**♦** (b) The average of all the integers

# Map task:

**Reduce task:** 

# Notes on example (b) (b) The average of all the integers

- ◆ In this case, don't worry whether the integers are unique
- **◆ Correction:** Map task: produce (key, (num of integers, sum of integers)) as output
  - ➤ Value is the (num, sum) tuple
  - > Key could be num or sum of integers or something else
- ◆ Single Reduce task: sum all sums of integers, sum num of integers from Reduce tasks, calculate average; emit (average, 1)
  - If multiple values for one key, means there happened to be overlap on the value of the key
  - > Extract values from tuples and add to sums
- ◆ Since each Map task summarizes large chunk of data by a single (key, (num, sum)) key-value pair, this should scale to use a **single Reducer** even with thousands of Map tasks. 6

# Exercise 2.3.1 (cont.)

Design MapReduce algorithms to take a very large file of integers and produce as output:

**♦** (c) The same set of integers, but with each integer appearing only once

Map task:

**Reduce task:** 

## Notes on example (c)

# (c) The same set of integers, but with each integer appearing only once

- ◆ Does not require sorting of final output
- ◆ Map task: emit (integer, 1) once only for each unique integer (e.g., maintain an array or list to track whether you have seen/emitted that integer before)
- **♦** Shuffle step will group together all values for the same integer: (integer, 1, 1, 1, 1...)
  - Same integer may appear in multiple chunks that go to different Map tasks
  - Each integer key will only go to one Reduce task
- **◆** Reduce task: eliminate duplicates (ignore list of 1's) for each integer key and emit (integer)
- ◆ Combine the output from multiple reduce tasks (not required to be in any order.

# Exercise 2.3.1 (cont.)

Design MapReduce algorithms to take a very large file of integers and produce as output:

**♦** (d) The <u>count</u> of the number of <u>distinct</u> integers in the input

Map task:

**Reduce task:** 

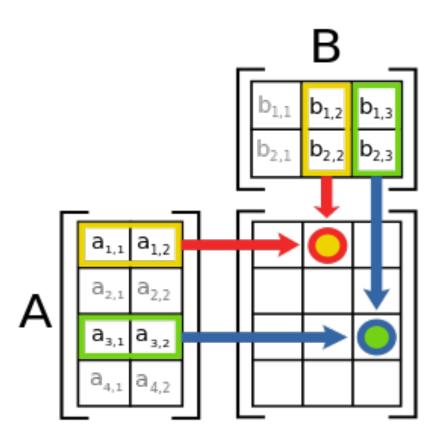
## **Intuition for (d)**

- (d) The count of the number of distinct integers in the input
  - ◆ Need to eliminate duplicates and then count unique integers
  - ◆ Each Map task only knows about its chunk of data
    - ➤ It can eliminate duplicates for that chunk (as in part (c))
  - ◆ Shuffle step: groups together any duplicates from multiple Map tasks before sending them to corresponding Reduce task
  - ◆ Reduce task can eliminate duplicates for the keys (integers) that it processes (as in part (c))
  - **◆** But can Reduce task(s) also count the unique integers?
    - > Can only count unique integers for the range of keys it processes
  - **◆** Can we assume a single Reduce task in this case?
    - ➤ No, because **input file is too large for a single machine**, and we don't know how many duplicates we will eliminate
  - ◆ If multiple Reduce tasks, then how do we count across? ¹º

# Notes on example (d)

- ◆ Two stages needed, since very large input file
- **◆** First phase: eliminate duplicates in large input file
  - ➤ Map task 1: just emit (integer, 1) for unique integers
  - > Sort by key, so disjoint set or range of integers goes to each reducer
  - Reduce task 1: Eliminates duplicates, counts unique integers, emits count of unique integers for its range of inputs (count, count)
  - Note that the Reduce tasks will process disjoint, non-overlapping sets of keys, so don't need to preserve the values
- **♦** Second phase: count unique numbers overall
  - ➤ Map task 2: Each map task will get some number of counts as input from the previous Reduce phase
  - > Adds them together and emits (count, count)
  - > Sort step: only one Reducer task, so all (count, count) pairs sent there
  - ➤ Reduce task 2: single Reducer, sums all counts from map tasks and produces overall count (could be multiple values for a key).

# **Matrix Multiplication**



# **Matrix:** Rectangular array of numbers:

4 rows
$$\overrightarrow{A} = \begin{bmatrix}
1402 & 191 \\
1371 & 821 \\
949 & 1437 \\
147 & 1448
\end{bmatrix}$$

$$\uparrow \qquad \uparrow \qquad \uparrow \qquad \uparrow \qquad \uparrow \qquad \uparrow \qquad 2 \text{ columns}$$



Dimension of matrix

: number of rows x number of columns

# **Matrix Elements** (entries of matrix)

$$A = \begin{bmatrix} 1402 & 191 \\ 1371 & 821 \\ 949 & 1437 \\ 147 & 1448 \end{bmatrix}$$

$$A12 = 191$$

$$A32 = 1437$$

$$A41 = 147$$

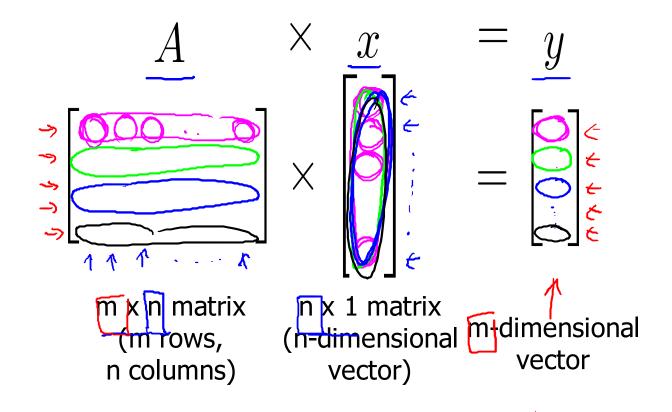
$$A43 = Undefined X$$

 $A_{ij} = "i,j$  entry" in the  $i^{th}$  row,  $j^{th}$  column.

# **Example: Matrix-Vector Multiplication**

$$\begin{bmatrix} 1 & 3 \\ 4 & 0 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 5 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 3 \end{bmatrix} \begin{bmatrix} 1 \\ 4 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ 5 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix}$$

# **Details:**



To get  $y_i$ , multiply A's  $i^{th}$ row with elements of vector x, and add them up.

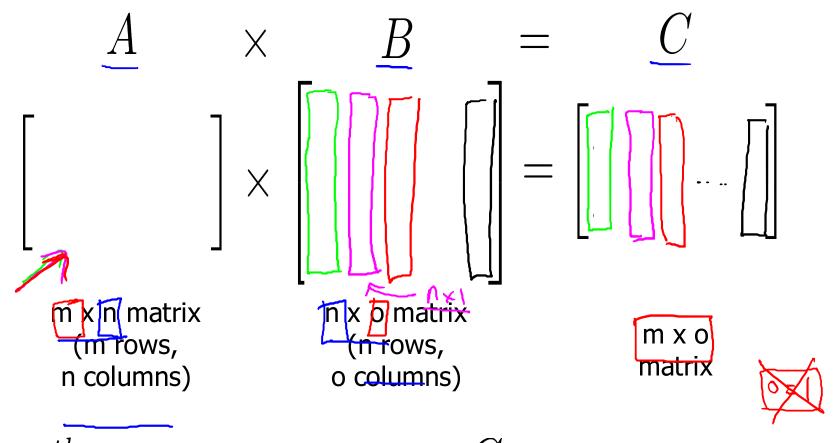
# **Example: Matrix Multiplication**

$$\begin{bmatrix} 1 & 3 & 2 \\ 4 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 3 \\ 0 & 1 \\ 5 & 2 \end{bmatrix} = \begin{bmatrix} 1 & 10 \\ 9 & 14 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 3 & 2 \\ 4 & 0 & 1 \end{bmatrix} \times \begin{bmatrix} 1 \\ 0 \\ 5 \end{bmatrix} = \begin{bmatrix} 11 \\ 9 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 3 & 2 \\ 4 & 0 & 1 \end{bmatrix} \times \begin{bmatrix} 3 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 10 \\ 14 \end{bmatrix}$$

# **Details:**



The  $i^{th}$  column of the matrix C is obtained by multiplying A with the  $i^{th}$  column of B. (for  $i=1,2,...,o)_{18}$ 

# **Example**

$$\begin{bmatrix} 1 & 3 \\ 2 & 5 \end{bmatrix} \begin{bmatrix} 0 \\ 3 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 9 & 7 \\ 15 & 12 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 3 \\ 2 & 5 \end{bmatrix} \begin{bmatrix} 0 \\ 3 \end{bmatrix} = \begin{bmatrix} 1 \times 0 + 3 \times 3 \\ 2 \times 0 + 5 \times 3 \end{bmatrix} = \begin{bmatrix} 9 \\ 15 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 3 \\ 2 & 5 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \times 1 + 3 \times 2 \\ 2 \times 1 + 5 \times 2 \end{bmatrix} = \begin{bmatrix} 7 \\ 12 \end{bmatrix}$$

# **Sparse matrix**

- **◆** Sparse matrix
  - A sparse matrix is a matrix in which most of the elements are zero.

$$\begin{bmatrix} 1.1 & 0 & 0 & 0 & 0 & 0 & 0.5 \\ 0 & 1.9 & 0 & 0 & 0 & 0 & 0.5 \\ 0 & 0 & 2.6 & 0 & 0 & 0 & 0.5 \\ 0 & 0 & 7.8 & 0.6 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1.5 & 2.7 & 0 & 0 \\ 1.6 & 0 & 0 & 0 & 0.4 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0.9 & 1.7 \end{bmatrix}$$

◆ Spars matrix representations (Row i, Column j, value)

# **Matrix Multiply Example**

- Matrix representations (Row i, Column j, value)
- Matrix Multiply:

$$C = A X B$$

A has dimensions L x M

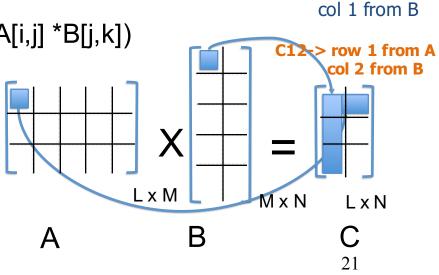
B has dimensions M x N

C has dimensions L x N

Matrix multiplication:  $C[i,k] = Sum_j (A[i,j] *B[j,k])$ 

Key = 
$$(i,k)$$
 for k in 1..N  
Value = A[i,j]

Key = 
$$(i,k)$$
 for i in 1..L  
Value =  $B[j,k]$ 



C11->row 1 from A

# **Matrix Multiply Example**



A has dimensions L x M

B has dimensions M x N

C has dimensions L x N

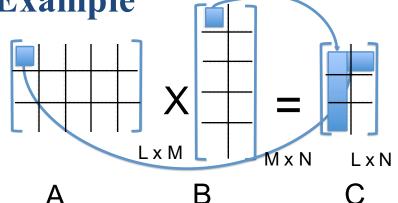
Matrix multiplication:  $C[i,k] = Sum_i (A[i,j] *B[j,k])$ 

### In the map phase:

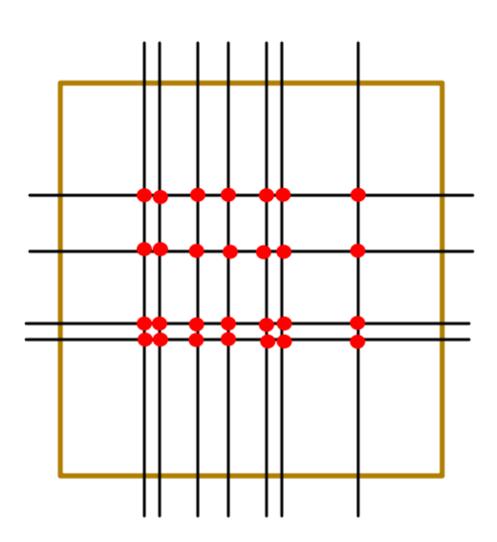
- for each element (i,j) of A, emit ((i,k), A[i,j]) for k in 1..N
  - Better: emit ((i,k)('A', i, k, A[i,j])) for k in 1..N
- for each element (j,k) of B, emit ((i,k), B[j,k]) for i in 1..L
  - Better: emit ((i,k)('B', i, k, B[j,k])) for i in 1..L

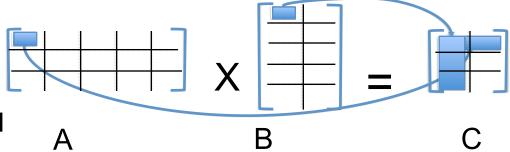
### In the reduce phase One reducer per output cell, emit

- key = (i,k)
- value = Sum<sub>j</sub> (A[i,j] \* B[j,k])



# The Responsibility of One Reducer





 $C[i,k] = Sum_j (A[i,j] *B[j,k]], C is LxN$ In the map phase:

- for each element (i,j) of A, emit ((i,k)('A', i, k, A[i,j])) for k in 1..N
- for each element (j,k) of B, : emit ((i,k)('B', i, k, B[j,k])) for i in 1..L

Example (what is needed at the reducer):

```
C[1,1] = A[1,1] * B[1,1] + A[1,2] * B[2,1] + A[1,3] * B[3,1] + A[1,4] * B[4,1] + A[1,5] * B[5,1]
```

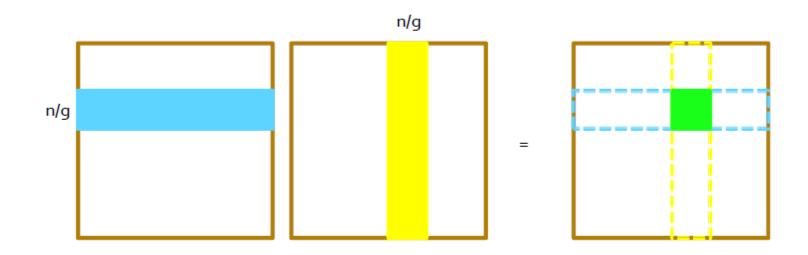
$$C[1,2] = A[1,1] * B[1,2] + A[1,2] * B[2,2] + A[1,3] * B[3,2] + A[1,4] * B[4,2] + A[1,5] * B[5,2]$$

$$C[2,1] = A[2,1] * B[1,1] + A[2,2] * B[2,1] + A[2,3] * B[3,1] + A[2,4] * B[4,1] + A[2,5] * B[5,1]$$

Map phase: For A[1,2], emit ('A', 1, k, A[1,2]) for k in 1..2 emit ((1,1)('A', 1, 1, A[1,2])), ((1,2)('A', 1, 2, A[1,2]))

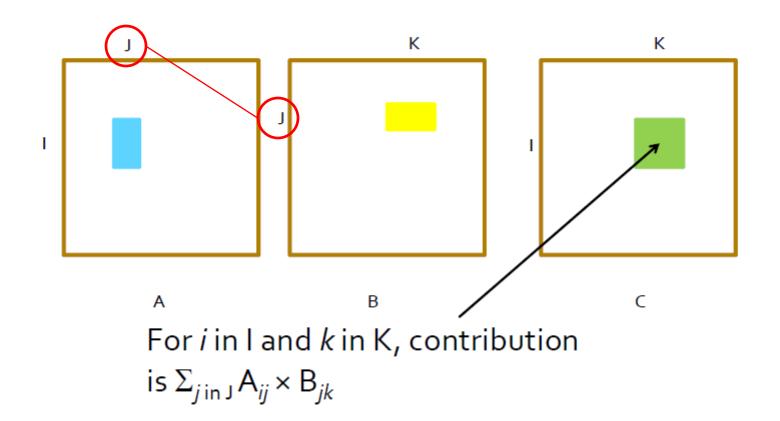
For **B[3,1]**: emit ('B', i, 1, B[3,1]) for i in 1..3 emit ((1,1)('B', 1, 1,B[3,1])), ((2,1)('B', 2, 1,B[3,1])), ((3,1)('B', 3, 1,B[3,1]))

# Picture of One Reducer



- A better way: use two map reduce jobs.
- ➤ Job 1: Divide both input matrices into rectangles.
  - Reducer takes two rectangles and produces partial sums of certain outputs.
- > Job 2: Sum the partial sums.

# **Picture of First Job**



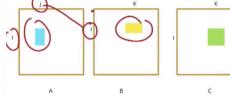
- Farhan Khwaja
- ▶ Idea: Multiply the appropriate values on the first MapReduce phase, Add on second MapReduce phase
- $ightharpoonup C[i,k] = Sum_j (A[i,j] *B[j,k])$
- C[1,1] = A[1,1] \* B[1,1] + A[1,2] \* B[2,1] + A[1,3] \* B[3,1] + A[1,4] \* B[4,1] + A[1,5] \* B[5,1]
- C[1,2] = A[1,1] \* B[1,2] + A[1,2] \* B[2,2] + A[1,3] \* B[3,2] + A[1,4] \* B[4,2] + A[1,5] \* B[5,2]Picture of First Job
- 1st Map Function:

For each matrix element A[ij]: emit(j, (A, i, A[i,j])

For each matrix element B[jk]: emit(j, (B, k, B[j,k])

Example (all will go to same reducer with key 2):

- For A[1,2]: emit (2, (A, 1, A[1,2])
- For B[2,1]: emit (2, (B, 1, B[2,1])
- For B[2,2]: emit (2, (B, 2, B[2,2])



### 1st Reduce Function:

Each key j will have its own Reduce function, will produce all products for a given value of j

```
For each value of (i,k) which comes from A and B, i.e. (A, i, A[ij]) and (B, k, B[jk]): emit((i,k), (A[ij] * B[jk]))
```

### **Example: Input from Map task**

```
For A[1,2]: emit (2, (A, 1, A[1,2])
```

For B[2,1]: emit (2, (B, 1, B[2,1])

For B[2,2]: emit (2, (B, 2, B[2,2])

### Reduce task for key j:

```
emit ( (1,1), A[1,2] * B[2,1])
emit ( (1,2), A[1,2] * B[2,2])
```

- C[1,1] = A[1,1] \* B[1,1] + A[1,2] \* B[2,1] + A[1,3] \* B[3,1] + A[1,4]
   \* B[4,1] + A[1,5] \* B[5,1]
   (Note: every term has key (1,1))
- $\circ$  C[1,2] = A[1,1] \* B[1,2] + A[1,2] \* B[2,2] + A[1,3] \* B[3,2] + A[1,4]<sub>29</sub> \* B[4,2] + A[1,5] \* B[5,2] (Note: every term has key (1,2)

- 2nd Map Function: map(key,value):
  #Let the pair of (((i,k), (A[ij] \* B[jk])) pass through
- 2nd Reduce Function:reduce(key,values):#each (i,k) will have its own reduce function
- For each (i,k) we will add up the values, thus emitting: emit((i,k),Sum(values))
- Example: C[1,1] = A[1,1] \* B[1,1] + A[1,2] \* B[2,1] + A[1,3] \* B[3,1] + A[1,4] \* B[4,1] + A[1,5] \* B[5,1]
- $\succ$  (Note: every term has key (1,1), will go to same reducer).  $\sqrt{\phantom{a}}$

# More Examples: Relational Join

### **Employee**

Name	SSN
Sue	99999999
Tony	77777777

### **Assigned Departments**

EmpSSN	DepName
99999999	Accounts
77777777	Sales
77777777	Marketing

## Emplyee ⋈ Assigned Departments

Name	SSN	EmpSSN	DepName
Sue	99999999	99999999	Accounts
Tony	77777777	77777777	Sales
Tony	77777777	77777777	Marketing

## Relational Join in MapReduce: Before Map Phase

### **Employee**

Name	SSN
Sue	99999999
Tony	77777777

### **Assigned Departments**

EmpSSN	DepName
99999999	Accounts
77777777	Sales
77777777	Marketing

Key idea: Lump all the tuples together into one dataset





### Relational Join in MapReduce: Map Phase



```
key=99999999, value=(Employee, Sue, 999999999)
key=77777777, value=(Employee, Tony, 77777777)
key=99999999, value=(Department, 999999999, Accounts)
key=77777777, value=(Department, 77777777, Sales)
key=77777777, value=(Department, 777777777, Marketing)
```

why do we use this as the key?

- Map Task: Emit (key, value) pair
- Key is unique ID in each table, used for join
- Value is a tuple with all fields from table (including key).

33

## Relational Join in MapReduce: Reduce Phase

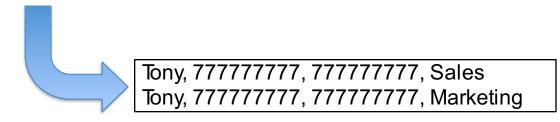
key=99999999, values=[(Employee, Sue, 99999999), (Department, 99999999, Accounts)]



Sue, 99999999, 99999999, Accounts

34

key=77777777, values=[(Employee, Tony, 777777777), (Department, 777777777, Sales), (Department, 77777777, Marketing)]



- Group by Key: groups together all values (tuples)
  associated with each key
- Reduce task: emit joined values (without table names).

# Relational Join in MapReduce, again

### Order(orderid, account, date)

### LineItem(orderid, itemid, qty)

1, aaa, d1 2, aaa, d2 3. bbb. d3

1, 10, 1 1, 20, 3 2, 10, 5 2, 50, 100 3. 20. 1

### Map Task

### relation name



- 1, aaa, d1
- 2, aaa, d2
- 3, bbb, d3

- → 1: "Order", (1,aaa,d1)
- → 2: "Order", (2,aaa,d2)
- → 3: "Order", (3,bbb,d3)

### Line

- 3, 20, 1

- 1, 10, 1  $\rightarrow$  1: "Line", (1, 10, 1)
- 1, 20, 3  $\rightarrow$  1: "Line", (1, 20, 3)
- $2, 10, 5 \rightarrow 2$ : "Line", (2, 10, 5)
- $2, 50, 100 \rightarrow 2$ : "Line", (2, 50, 100)
  - → 3: "Line", (3, 20, 1)

### Reducer for key 1

"Order", (1,aaa,d1)

"Line", (1, 10, 1)

"Line", (1, 20, 3)



(1, aaa, d1, 1, 10, 1)

(1, aaa, d1, 1, 20, 3)

### **Distributed Sort**

➤ Sort a very large list of (firstName, lastName) pairs by lastName followed by firstName

- ➤ Map task:?
- ➤ Reduce task:?

# Examples of outputs:

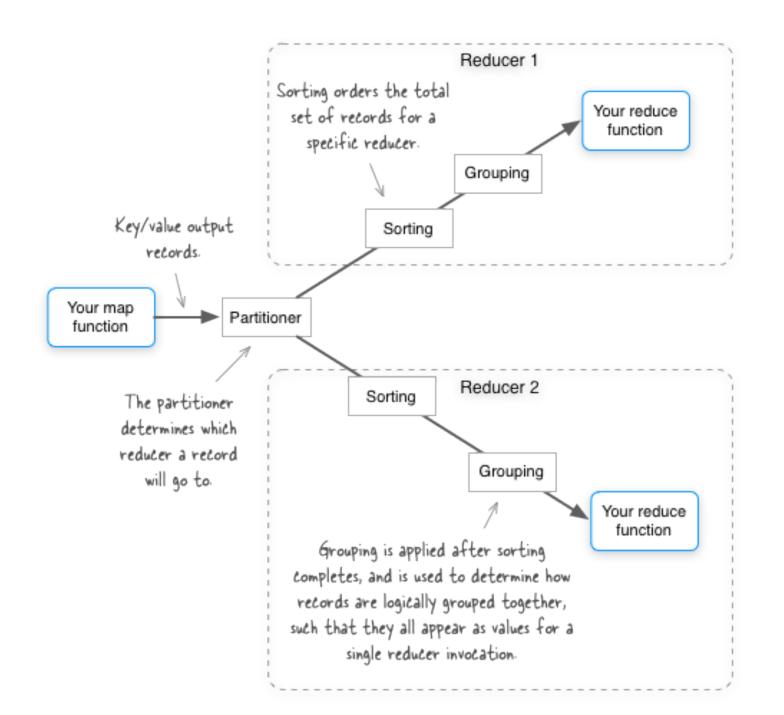
Smith Anne

Smith John

Smith Ken

## Distributed Sort explanation Sort by lastName followed by firstName

- Map task
  - ➤ Emit (lastName, firstName)
- o Group by keys:
  - > Group together entries with same last name
  - ➤ Divide into non-overlapping alphabetical ranges
  - > Keys are sorted in alphabetical order
- Reduce task
  - > Processes one key at a time (in alphabetical order)
  - ➤ For each lastName key, if there are multiple firstName values, will emit (lastName, firstName) in alphabetical order
  - ➤ Merge output from all Reduce tasks.



## **Example: Join By Map-Reduce**

- Compute the natural join  $R(A,B) \bowtie S(B,C)$
- R and S are each stored in files
- Tuples are pairs (a,b) or (b,c)

Α	В
a <sub>1</sub>	b <sub>1</sub>
$a_2$	b <sub>1</sub>
$a_3$	$b_2$
$a_4$	$b_3$



В	C	
$b_2$	C <sub>1</sub>	
$b_2$	<b>C</b> <sub>2</sub>	=
$b_3$	<b>C</b> <sub>3</sub>	

Α	С
$\mathbf{a}_3$	<b>C</b> <sub>1</sub>
$a_3$	<b>C</b> <sub>2</sub>
$a_4$	<b>C</b> <sub>3</sub>

S

R

## Map-Reduce Join

- Use a hash function h from B-values to 1...k
- A Map process turns:
  - Each input tuple R(a,b) into key-value pair (b,(a,R))
  - Each input tuple S(b,c) into (b,(c,S))
- Map processes send each key-value pair with key b to Reduce process h(b)
  - Hadoop does this automatically; just tell it what k is.
- Each Reduce process matches all the pairs (b,(a,R)) with all (b,(c,S)) and outputs (a,b,c).

## **Cost Measures for Algorithms**

- In MapReduce we quantify the cost of an algorithm using
- Communication cost = total I/O of all processes
- 2. Elapsed communication cost = max of I/O along any path
- 3. (*Elapsed*) *computation cost* analogous, but count only **running time of processes**.

Note that here the big-O notation is not the most useful (adding more machines is always an option)

## **Example: Cost Measures**

- For a map-reduce algorithm:
  - Communication cost =
    - input file size +
    - sum of the sizes of all files passed from Map processes to Reduce processes +
    - the sum of the output sizes of the Reduce processes.
  - Elapsed communication cost is the sum of the largest input + output for any map process, plus the same for any reduce process.

#### **What Cost Measures Mean**

- Either the I/O (communication) or processing (computation) cost dominates
  - Ignore one or the other
- Total cost tells what you pay in rent from your friendly neighborhood cloud
- Elapsed cost is wall-clock time using parallelism.

## **Cost of Map-Reduce Join**

- Total communication cost
  - $= O(|R|+|S|+|R\bowtie S|)$
- Elapsed communication cost = O(s)
  - Choose the number of Map processes so that the I/O limit s is respected
  - s is the amount of input or output that any one process can have
  - s could be:
    - What fits in main memory
    - What fits on local disk
- With proper indexes, computation cost is linear in the input + output size
  - So computation cost is like comm. Cost.

# Problems Suited for MapReduce

Data set is truly "big"

- o Terabytes, not tens of gigabytes
- Hadoop/MapReduce designed for terabyte/petabyte scale computation
- Most real-world problems process less than 100 Gbytes of input
  - ➤ Microsoft, Yahoo: median job under 14 GB
  - Facebook: 90% of jobs under 100 GB.

#### Don't need fast response time

- When submitting jobs, Hadoop latency can be a minute
- Not well-suited for problems that require faster response time
  - online purchases, transaction processing
- A good pre-computation engine
  - ➤ E.g., pre-compute related items for every item in inventory.

Good for applications that work in batch mode

- o Jobs run without human interaction, intervention
- o Runs over entire data set
  - Takes time to initiate, run; shuffle step can be timeconsuming
- Does not support real time applications well: sensor data, real-time advertisements, etc.
- Does not provide good support for random access to datasets
  - Extensions: Hive, Dremel, Shark, Amplab.

- ◆ Best suited for data that can be expressed as keyvalue pairs without losing context, dependencies
- Graph data harder to process using MapReduce
- Implicit relationships: edges, sub-trees, child/parent relationships, weights, etc.
- Graph algorithms may need information about the entire graph for each iteration
  - ➤ Hard to break into independent chunks for Map tasks
- Alternatives: Google's Pregel, Apache Giraph.

Other problems/data \*not\* suited for MapReduce

- Tasks that need results of intermediate steps to compute results of current step
  - > Interdependencies among tasks
  - ➤ Map tasks must be independent
- Some machine learning algorithms
  - ➤ Gradient-based learning, expectation maximization.

### Problems NOT suitable for MapReduce

- MapReduce is great for:
  - Problems that require sequential data access
  - Large batch jobs (not interactive, real-time)
- MapReduce is inefficient for problems where random (or irregular) access to data required:
  - Graphs
  - Interdependent data
    - Machine learning
    - Comparisons of many pairs of items.

#### **Summary:**

#### **Good candidates for MapReduce:**

- Jobs that have to process huge quantities of data and either summarize or transform the contents
- Collected data has elements that can easily be captured with an identifier (key) and corresponding value.

#### MapReduce inside Google



Googlers' hammer for 80% of our data crunching

- Large-scale web search indexing
- Clustering problems for <u>Google News</u>
- Produce reports for popular queries, e.g. <u>Google Trend</u>
- Processing of <u>satellite imagery data</u>
- Language model processing for <u>statistical machine</u> translation
- Large-scale <u>machine learning problems</u>
- Just a plain tool to reliably spawn large number of tasks o e.g. parallel data backup and restore

The other 20%? e.g. Pregel

































