

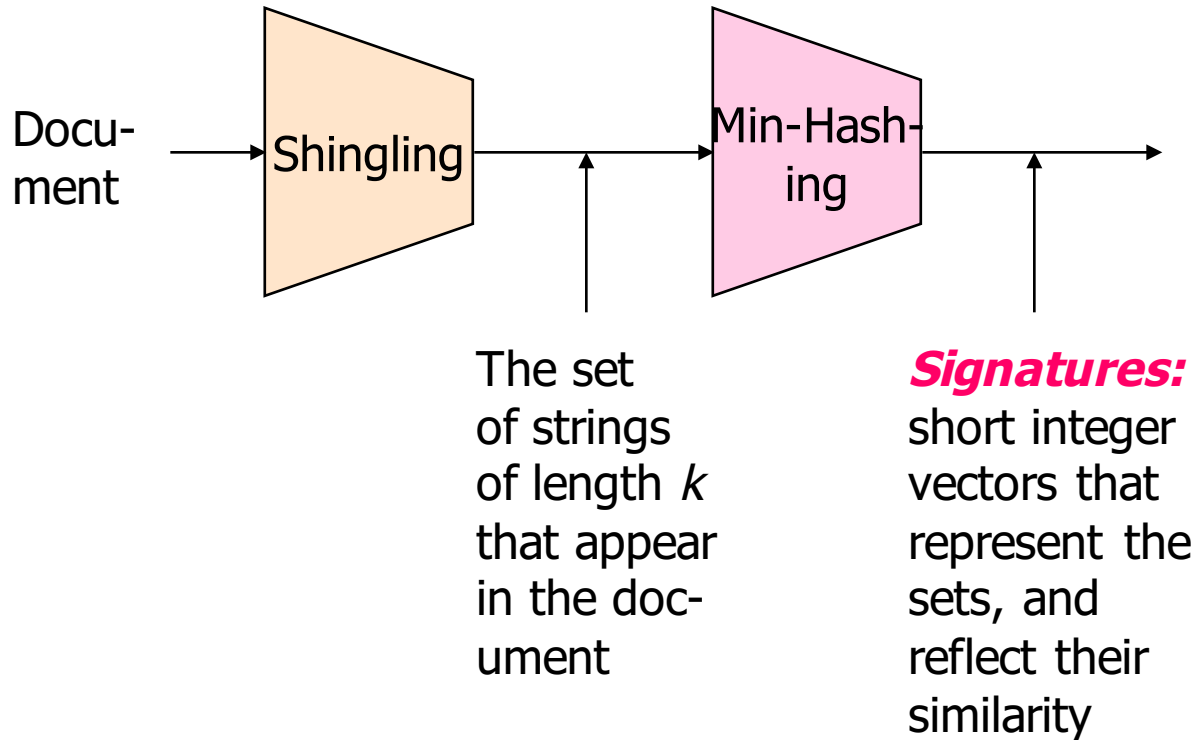
Finding Similar Sets (part 2)

Applications

Shingling

Minhashing

Locality-Sensitive Hashing



MinHashing

Step 2: *Minhashing*: Convert large sets to short signatures, while preserving similarity

MinHashing

Data as Sparse Matrices

Jaccard Similarity Measure

Constructing Signatures

From Sets to Boolean Matrices

- ◆ **Rows** = **elements** of the universal set
- ◆ **Columns** = **sets**
- ◆ **1** in **row e** and **column S** if and only if **element e is a member of set S**
- ◆ Column similarity is the **Jaccard similarity** of the sets of their rows with 1: **intersection/union of sets**
- ◆ **Typical matrix is sparse** (many 0 values)
 - May not really represent the data by a boolean matrix
 - Sparse matrices are usually better represented by the list of non-zero values (e.g., triples)
 - But the matrix picture is conceptually useful.

Example 3.6

<i>Element</i>	S_1	S_2	S_3	S_4
a	1	0	0	1
b	0	0	1	0
c	0	1	0	1
d	1	0	1	1
e	0	0	1	0

- ◆ Universal set: $\{a, b, c, d, e\}$
- ◆ Matrix represents sets chosen from universal set
- ◆ $S_1 = \{a, d\}$, $S_2 = \{c\}$, $S_3 = \{b, d, e\}$ and $S_4 = \{a, c, d\}$
- ◆ Example: **rows are products** and **columns are customers**, represented by **set of items** they bought
- ◆ **Jacquard similarity of S_1 , S_4** : intersection/union = $2/3$.

Example: Jaccard Similarity of Columns

C₁—C₂

0 1 *

1 0 *

1 1 * *

0 0

1 1 * *

0 1 *

$$\text{Sim}(C_1, C_2) = \\ 2/5 = 0.4$$

Outline: Finding Similar Columns

1. Compute **signatures** of columns = **small summaries** of columns
2. Examine **pairs of signatures** to find similar signatures
 - **Essential:** similarities of signatures and columns are related
3. **Optional:** check that columns with similar signatures are really similar.

Warnings

1. Comparing **all pairs of signatures** may take **too much time**, even if not too much space
 - A job for **Locality-Sensitive Hashing**
2. These methods can produce false negatives, and even false positives (if the optional check is not made).

Signatures

- ◆ Key idea: “hash” each column C to a small *signature* $Sig(C)$, such that:
 1. $Sig(C)$ is **small** enough that we can fit a signature in **main memory** for each column
 2. $Sim(C_1, C_2)$ is the same as the “**similarity**” of $Sig(C_1)$ and $Sig(C_2)$.

Four Types of Rows

- ◆ Given columns C_1 and C_2 , rows may be classified as:

	<u>C_1</u>	<u>C_2</u>
a	1	1
b	1	0
c	0	1
d	0	0

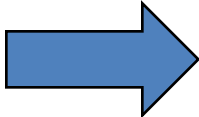
- ◆ Also, a = # rows of type a , etc.
- ◆ Note $Sim(C_1, C_2) = a / (a + b + c)$
 - Jacquard similarity: intersection/union
 - a is intersection, $a+b+c$ is union

Minhashing

1. To *minhash* a set represented by a column of the matrix, **pick a random permutation of the rows**
2. Define “hash” function $h(C)$ = the number of the first (in the permuted order) row in which column C has 1
3. Use several (e.g., 100) independent hash functions to **create a signature**.

Minhashing Example (3.7)





<i>Element</i>	S_1	S_2	S_3	S_4
<i>a</i>	1	0	0	1
<i>b</i>	0	0	1	0
<i>c</i>	0	1	0	1
<i>d</i>	1	0	1	1
<i>e</i>	0	0	1	0



 Permute

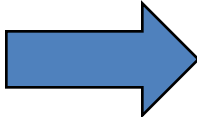
<i>Element</i>	S_1	S_2	S_3	S_4
<i>b</i>	0	0	1	0
<i>e</i>	0	0	1	0
<i>a</i>	1	0	0	1
<i>d</i>	1	0	1	1
<i>c</i>	0	1	0	1

1. To minhash a set represented by a column of the characteristic matrix, pick a **permutation of the rows**
2. The **minhash value** of any column is the number of first row, in permuted order, in which column **has a 1**
3. For set S_1 , first 1 appears in row *a*, so:

- $h(S_1) =$ 
- $h(S_2) =$ 
- $h(S_3) =$ 
- $h(S_4) =$ 

Minhashing Example (3.7)

<i>Element</i>	S_1	S_2	S_3	S_4
<i>a</i>	1	0	0	1
<i>b</i>	0	0	1	0
<i>c</i>	0	1	0	1
<i>d</i>	1	0	1	1
<i>e</i>	0	0	1	0

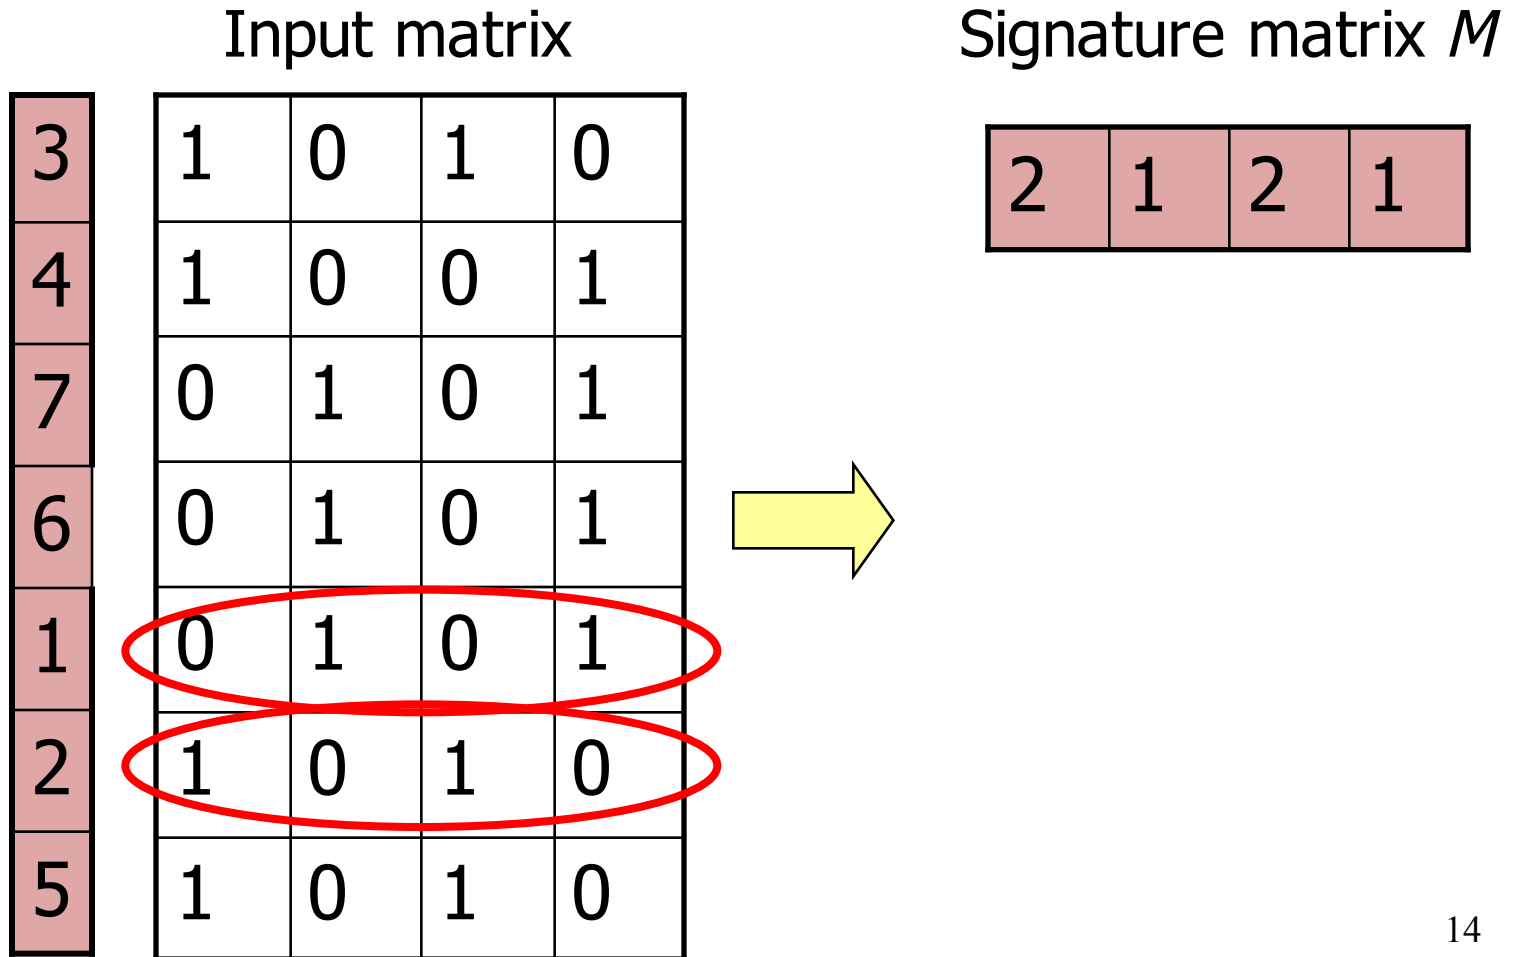


Permute

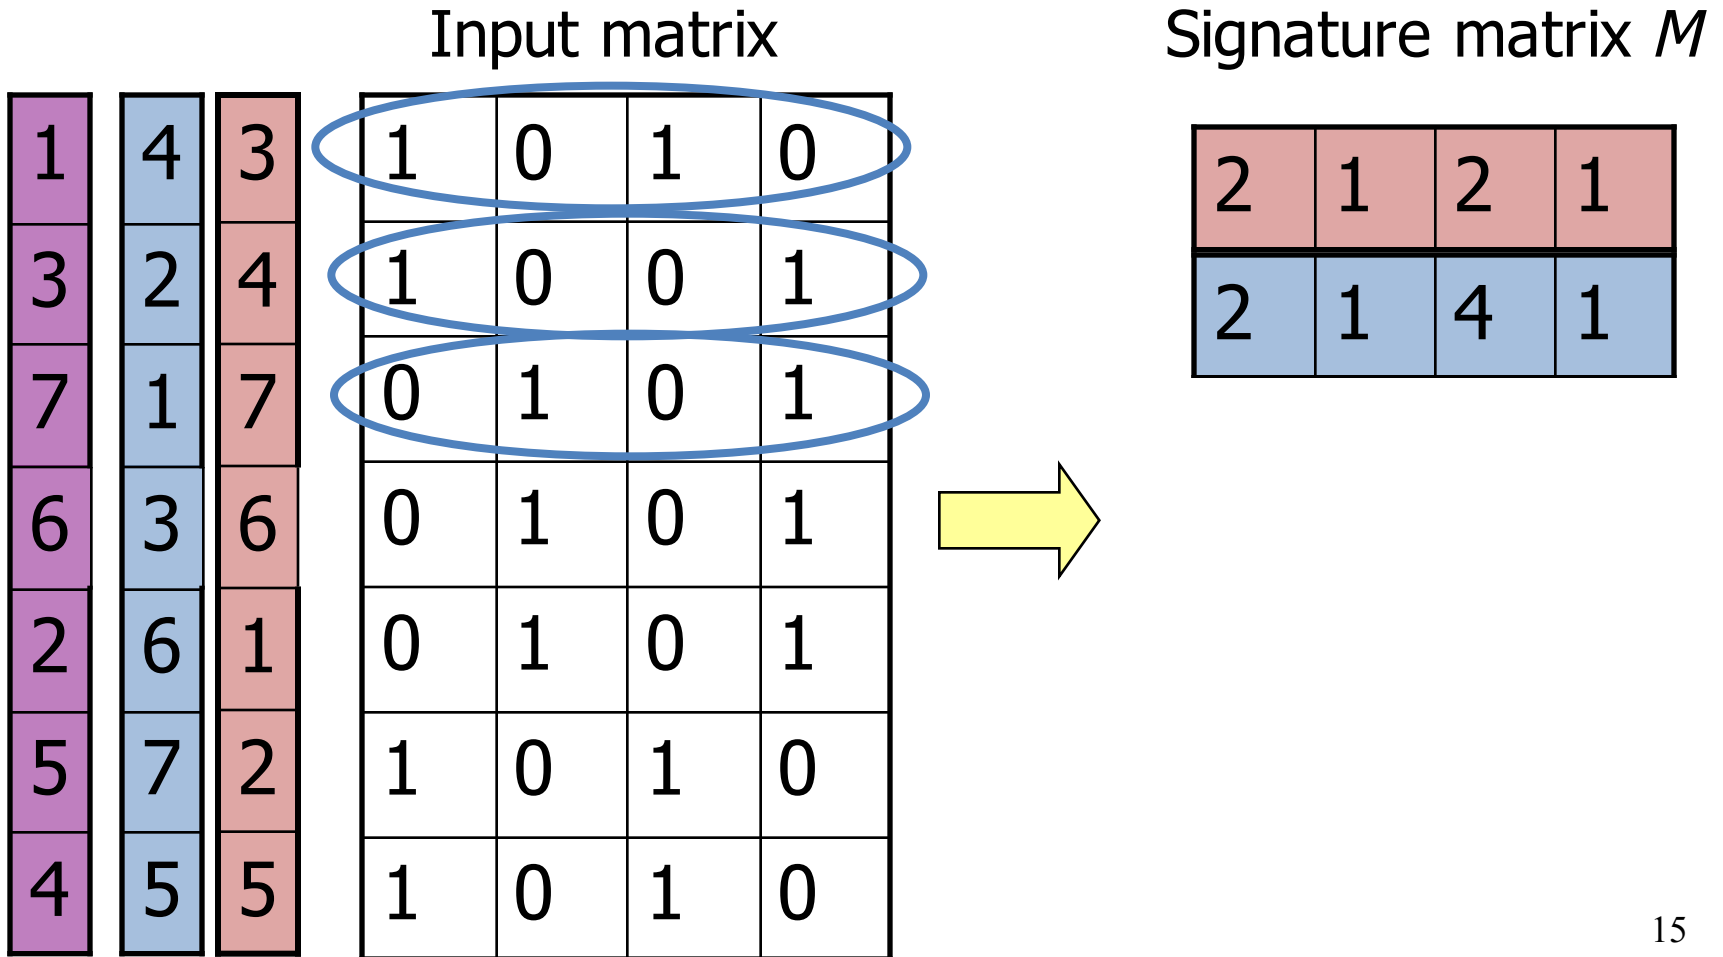
<i>Element</i>	S_1	S_2	S_3	S_4
<i>b</i>	0	0	1	0
<i>e</i>	0	0	1	0
<i>a</i>	1	0	0	1
<i>d</i>	1	0	1	1
<i>c</i>	0	1	0	1

1. To minhash a set represented by a column of the characteristic matrix, pick a **permutation of the rows**
2. The **minhash value** of any column is the number of first row, in permuted order, in which column **has a 1**
3. For set S_1 , first 1 appears in row *a*, so:
 - $h(S_1) = a$
 - $h(S_2) = c$
 - $h(S_3) = b$
 - $h(S_4) = a$

Minhashing Example



Minhashing Example




Minhashing Example

1	4	3
3	2	4
7	1	7
6	3	6
2	6	1
5	7	2
4	5	5

Input matrix

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0



Signature matrix M

2	1	2	1
2	1	4	1
1	2	1	2

Surprising Property: Connection between Minhashing and Jaccard Similarity

- ◆ The probability that minhash function for a **random permutation of rows** produces same value for two sets equals **Jaccard similarity** of those sets
 - Probability that $h(C_1) = h(C_2)$ is the same as $Sim(C_1, C_2)$
- ◆ Recall four types of rows:


	C_1	C_2
a	1	1
b	1	0
c	0	1
d	0	0

- ◆ $Sim(C_1, C_2)$ for both Jacquard and Minhash are $a / (a + b + c)$!
 - **Why?** Look down the permuted columns C_1 and C_2 until we see a 1
 - If it's a type-*a* row, then $h(C_1) = h(C_2)$. If a type-*b* or type-*c* row, then not. (Don't count the *type-d* rows).

Similarity for Signatures

- ◆ **Sets represented** by characteristic **matrix M**
- ◆ **To represent sets:** pick at random some number **n** of **permutations** of the rows of M
 - **100 permutations** or several hundred
- ◆ Call **minhash** functions determined by these permutations h_1, h_2, \dots, h_n
- ◆ From **column representing set S**, construct **minhash signature for S**:
 - vector $[h_1(S), h_2(S), \dots, h_n(S)]$, usually represented as column
- ◆ Construct a **signature matrix**: **i^{th} column of M** replaced by **minhash signature for i^{th} column**
- ◆ The **similarity of signatures** is the fraction of the hash functions in which they agree.

Min Hashing – Example




Min Hashing – Example

1	4	3
3	2	4
7	1	7
6	3	6
2	6	1
5	7	2
4	5	5

Input matrix

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0



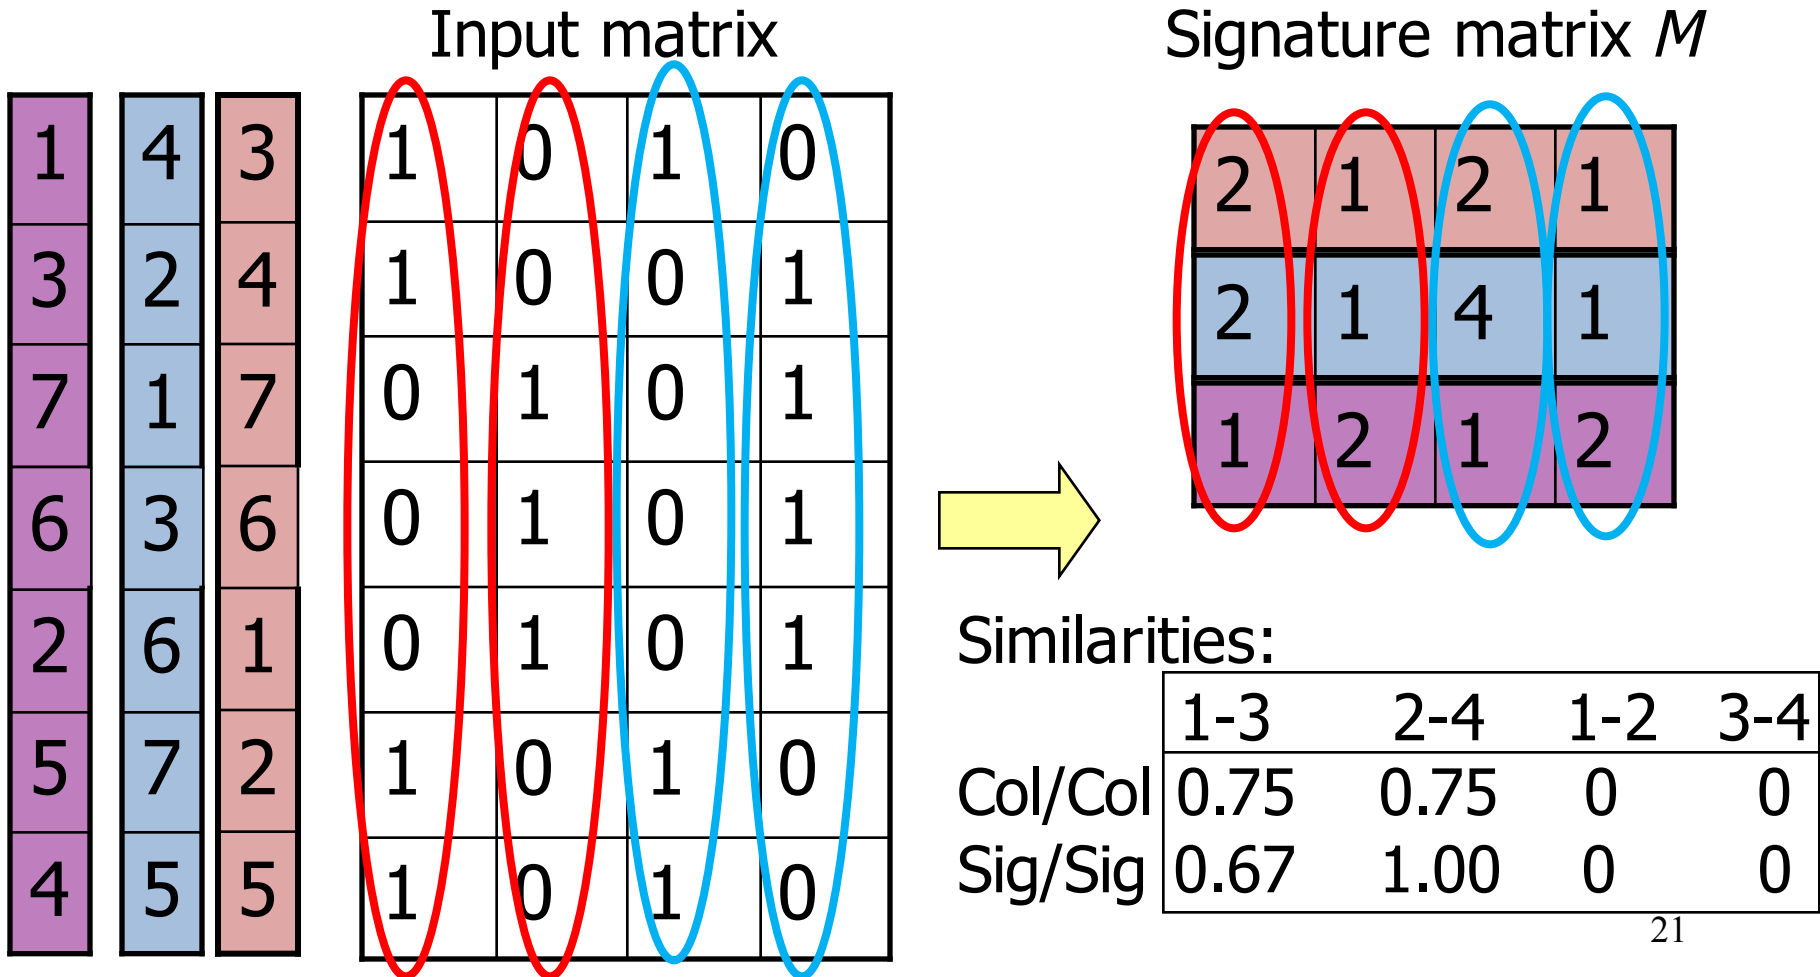
Signature matrix M

2	1	2	1
2	1	4	1
1	2	1	2

Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.67	1.00	0	0

Min Hashing – Example



Minhash Signatures

- ◆ Pick (say) **100 random** permutations of the rows
- ◆ Think of ***Sig* (C)** as a column vector
- ◆ Let ***Sig* (C)[i]** =
according to the ***i* th permutation**, the number of the first row that has a **1 in column C**.

Implementation – (1)

- ◆ **Not feasible to permute** a large characteristic matrix explicitly
 - Suppose **1 billion rows**
 - Hard to pick a **random permutation from 1...billion**
 - Representing a random permutation **requires 1 billion entries**
 - Accessing rows in permuted order leads to thrashing
- ◆ Can simulate the effect of a random permutation by a **random hash function**
 - **Maps row** numbers to as many buckets as there are rows
 - May have **collisions on buckets**
 - **Not important as long as number of buckets is large.**

Implementation – (2)

- ◆ A good approximation to permuting rows:
pick around 100 hash functions
- ◆ For each:
 - column c (set representing a document)
 - hash function h_i
- ◆ Keep a “slot” in signature matrix $M(i, c)$
- ◆ **Intent:** $M(i, c)$ will become the smallest value of $h_i(r)$ for which column c has 1 in row r
 - $h_i(r)$ gives order of rows for i^{th} permutation.

Implementation – (3)

```
for each row  $r$   
  for each column  $c$   
    if  $c$  has 1 in row  $r$   
      for each hash function  $h_i$  do  
        if  $h_i(r)$  is a smaller value than  $M(i, c)$  then  
           $M(i, c) := h_i(r);$ 
```

Computing Minhash Signatures:

Example 3.8

Row	S_1	S_2	S_3	S_4	$x + 1 \mod 5$	$3x + 1 \mod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

Two hash functions give permutations of rows:
 $h1 = x+1 \mod 5$, $h2 = 3x + 1 \mod 5$

	S_1	S_2	S_3	S_4
h_1	∞	∞	∞	∞
h_2	∞	∞	∞	∞

	S_1	S_2	S_3	S_4
h_1	1	∞	∞	1
h_2	<u>1</u>	∞	∞	<u>1</u>

Initial signature matrix

For row 0: Replace existing signature values with lower hash values for S_1 and S_4 , since both have 1 in row

Computing Minhash Signatures: Example 3.8 (part 2)

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	S_1	S_2	S_3	S_4
h_1	1	∞	<u>2</u>	1
h_2	1	∞	<u>4</u>	1

For row 1: replace h_1 and h_2 values for S_3 , since row has a 1 and values are lower

	S_1	S_2	S_3	S_4
h_1	1	<u>3</u>	2	1
h_2	1	<u>2</u>	4	1

For row 2: replace values for S_2 since set has a 1 value. Do not replace values for S_4 , because existing values are lower

Computing Minhash Signatures: Example 3.8 (part 3)

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3


	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	<u>0</u>	2	<u>0</u>	<u>0</u>

For row 3: don't replace h_1 values--all are below 4; replace h_2 values with 0 for S_1 , S_3 , S_4

	S_1	S_2	S_3	S_4
h_1	1	3	<u>0</u>	1
h_2	0	2	<u>0</u>	0

For row 4: replace h_1 value for S_3 , don't replace h_2 value since current value is lower

Note: result is same as ²⁸
permutations to find first 1



Locality Sensitive Hashing

Step 3: **Locality-Sensitive Hashing:**

Focus on pairs of signatures likely to be from similar documents

Example

Row	C1	C2
1	1	0
2	0	1
3	1	1
4	1	0
5	0	1

$$h(x) = x \bmod 5$$

$$g(x) = 2x+1 \bmod 5$$

	Sig1	Sig2
$h(1) = 1$	1	-
$g(1) = 3$	3	-
$h(2) = 2$	1	2
$g(2) = 0$	3	0
$h(3) = 3$	1	2
$g(3) = 2$	2	0
$h(4) = 4$	1	2
$g(4) = 4$	2	0
$h(5) = 0$	1	0
$g(5) = 1$	2	0