

Clustering

(part 2)

INF 553:
Foundations and Applications of Data Mining

Clustering

◆ **Clustering:** Given a **set of points**, with a notion of **distance** between points, **group the points** into some number of *clusters*

◆ **Algorithms:**

- Agglomerative **hierarchical clustering**:
 - Centroid (for Euclidean spaces)
 - Clustroid (for non-Euclidean spaces)

◆ **Point assignment**

- K-means
- BFR: extend k-means to handle large data set
- CURE.

***k*-means clustering**

k -means Algorithm(s)

◆ Point Assignment Algorithm

◆ Assumes Euclidean space/distance

◆ Start by picking k , the number of clusters


◆ Initialize clusters by picking one point per cluster

➤ **Example:** Pick one point at random, then $k-1$ other points, each as far away as possible from the previous points.

Populating Clusters

- ◆ 1) For each point, **place it in the cluster whose current centroid it is nearest**
- ◆ 2) **After all points are assigned, update the locations of centroids of the k clusters**
- ◆ 3) **Reassign all points to their closest centroid**
 - Sometimes moves points between clusters
- ◆ **Repeat 2 and 3 until convergence**
 - **Convergence:** Points don't move between clusters and centroids stabilize.

Example: Assigning Clusters




x ... data point

□ ... centroid

Clusters after round 1


Example: Assigning Clusters



x ... data point
□ ... centroid

Clusters after round 2

Example: Assigning Clusters




x ... data point
□ ... centroid

Clusters at the end

Getting the k right

How to select k ?

- ◆ Try different k , looking at the change in the average distance to centroid as k increases
- ◆ Average falls rapidly until right k , then changes little





Intuition on Picking k

- ◆ Consider a data set with three natural clusters
- ◆ If we force a merge of two of these natural clusters, the diameter jumps quickly
 - **Diameter of cluster** = **maximum** distance between any **two points** of the cluster
 - **Radius of cluster** = **maximum** distance between all **points and the centroid**
 - Measure of **appropriateness for a cluster** (e.g., average radius or diameter) **grows slowly**, as long as the number of clusters we assume is at or above the true number of clusters.



Example: Picking k

Too few;
many long
distances
to centroid.




Example: Picking k


Just right;
distances
rather short.



Example: Picking k



Too many;
little improvement
in average
distance.



Picking the initial K points

◆ Approach 1: Sampling

- **Cluster a sample** of the data using **hierarchical**
- clustering, to obtain k clusters
- Pick a point from each cluster (e.g. point closest to centroid)
- Sample fits in main memory

◆ Approach 2: Pick “dispersed” set of points

- Pick first point at random
- Pick the next point to be the one whose minimum distance from the selected points is **as large as possible**
- Repeat until we have k points.

Complexity


- ◆ In each round, we have to examine each **input point exactly once** to find **closest centroid**
- ◆ Each round is $O(kN)$ for **N points, k clusters**
- ◆ **But the number of rounds to convergence can be very large!**
- ◆ Can we cluster in a single pass over the data?

The BFR Algorithm


Extension of *k*-means to large data

BFR Algorithm

- ◆ **BFR** [Bradley-Fayyad-Reina]
- ◆ **Extension of k -means to large data**
- ◆ is a variant of k -means designed to handle **very large** (disk-resident) data sets
- ◆ **Assumes** that clusters are **normally distributed** around a centroid in a Euclidean space
 - Standard deviations in different dimensions may vary
 - Clusters are axis-aligned ellipses
- ◆ **Efficient way to summarize clusters**
(want memory required $O(\text{clusters})$ and not $O(\text{data})$).




Normal Distribution



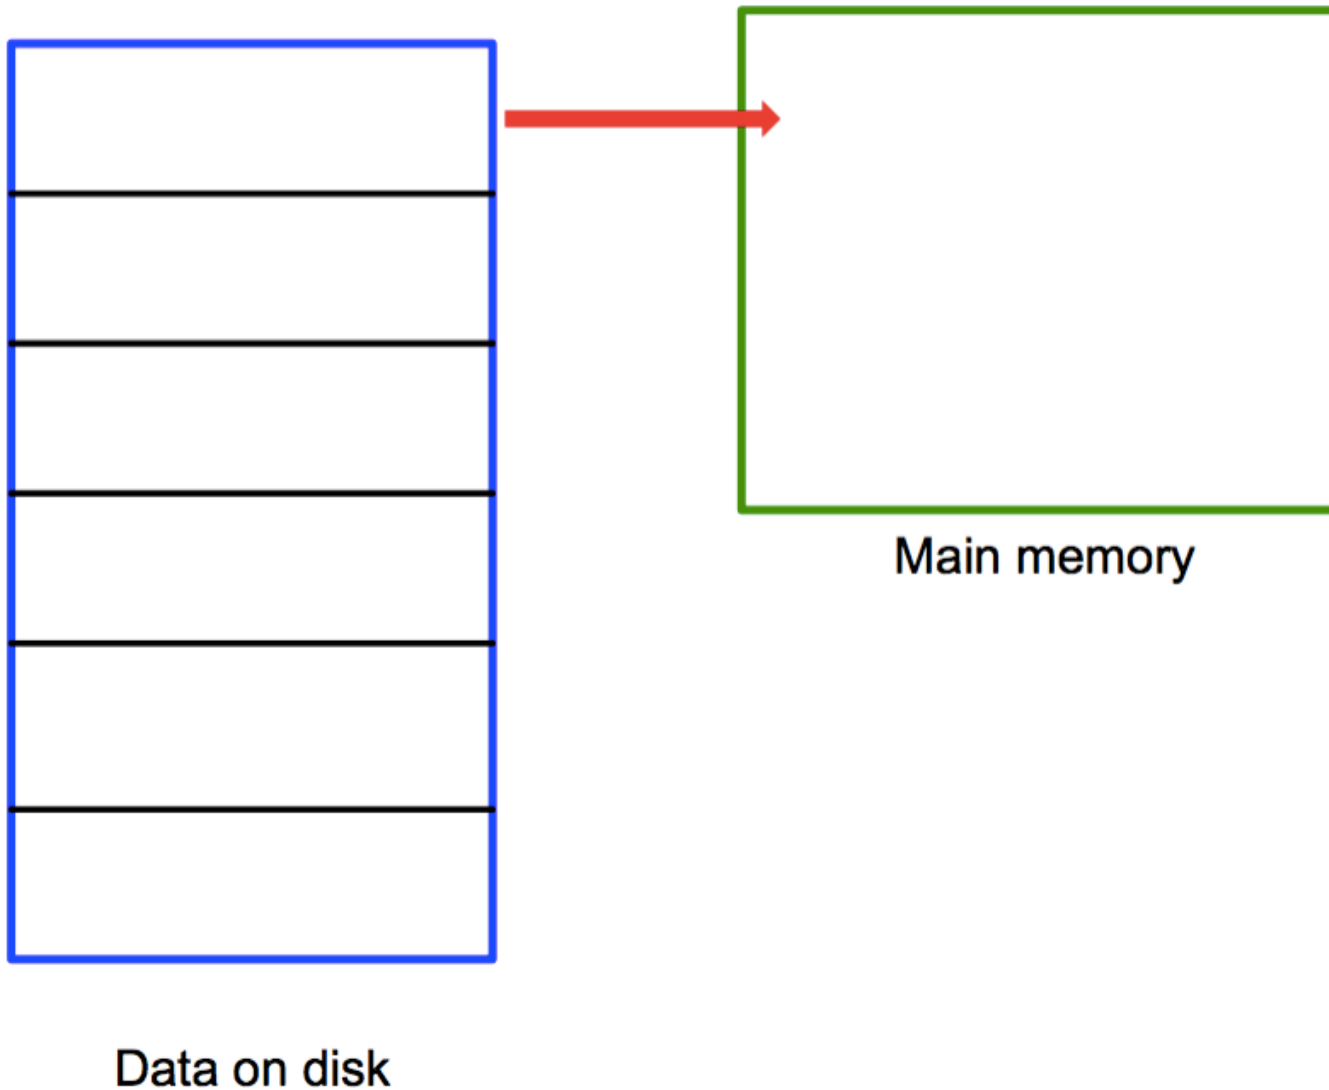
- ◆ Can quantify the likelihood of finding a point in the cluster at a given distance from the centroid along each dimension
- ◆ Standard deviation in different dimensions may vary

BFR Clusters

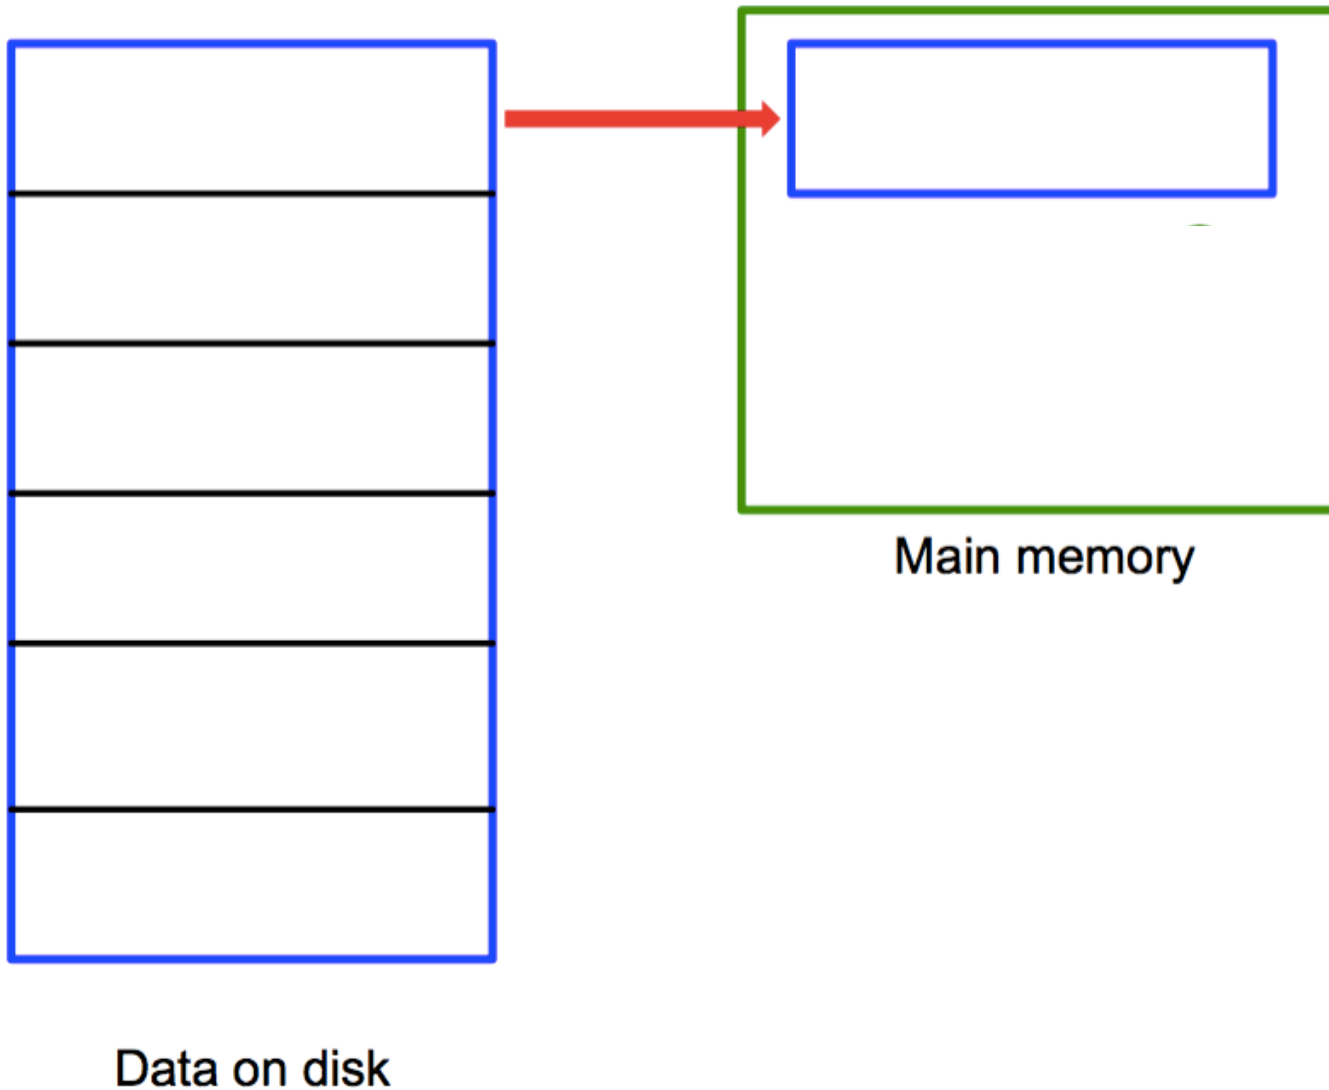
- ◆ Normal distribution assumption implies that clusters “look like” axis-aligned ellipses



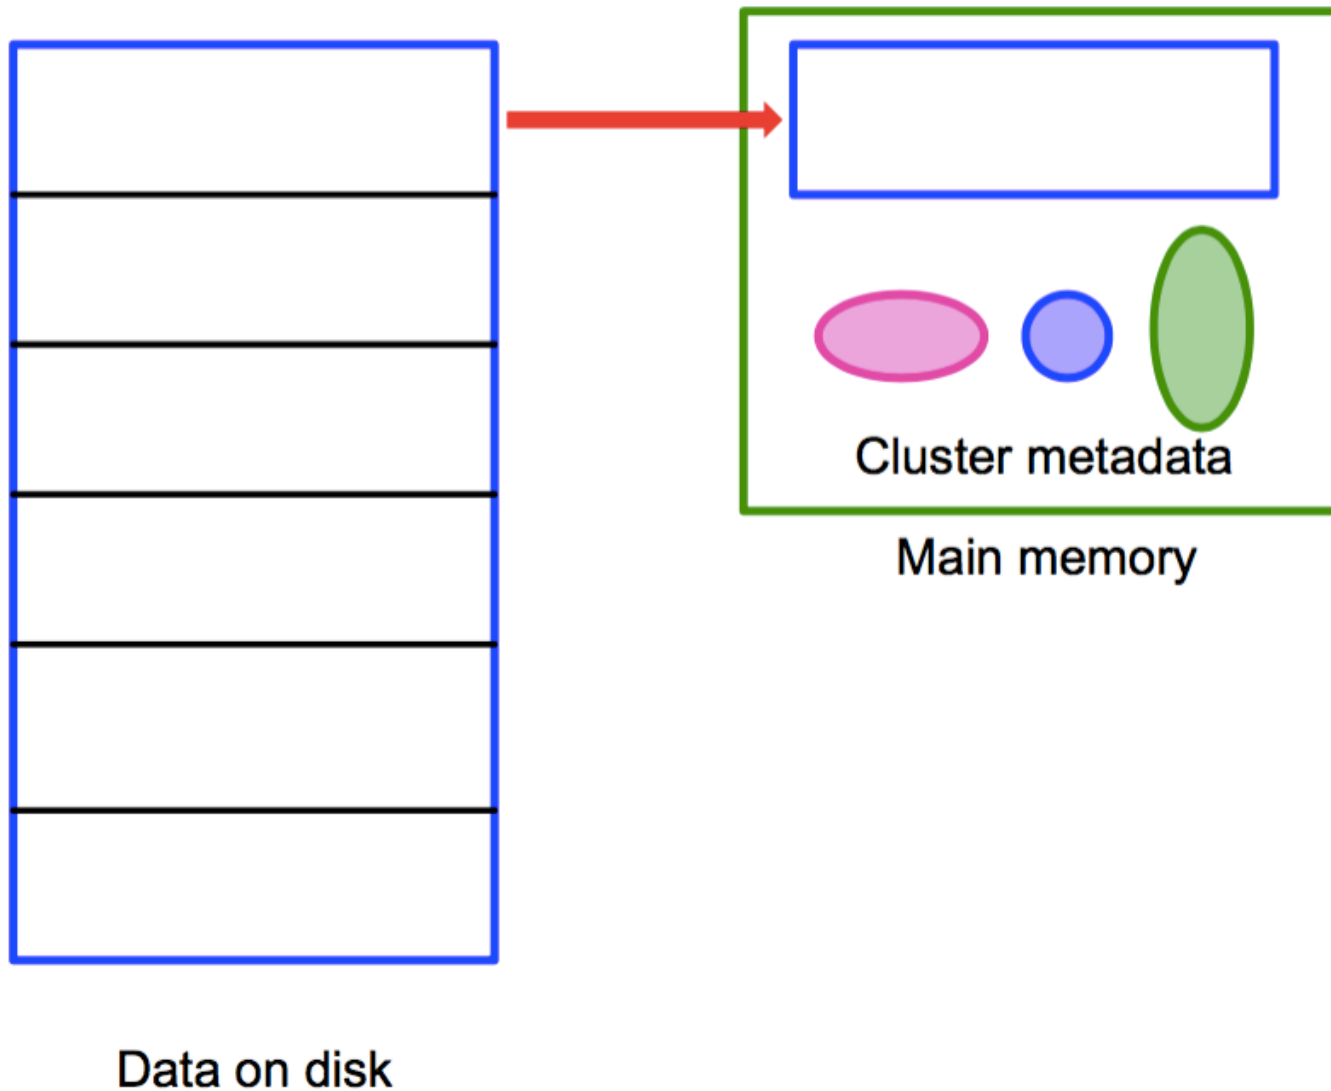
BFR Algorithm: Overview



BFR Algorithm: Overview



BFR Algorithm: Overview



BFR Algorithm

- ◆ Points are read from disk one main-memory-full at a time
- ◆ Most points from previous memory loads are summarized by **simple statistics**
- ◆ To begin, from the initial load we select the initial k centroids by some sensible approach:
 - Take k random points
 - Take a small random sample and cluster optimally
 - Take a sample; pick a random point, and then $k-1$ more points, each as far from the previously selected points as possible.

Three Classes of Points

3 sets of points which we keep track of:

◆ Discard set (DS):

- Points close enough to a centroid to be summarized


◆ Compression set (CS):

- Groups of points that are **close together** but not close to any existing centroid
- These points are summarized, but not assigned to a cluster
- Called a “miniclust”

◆ Retained set (RS):

- Isolated points waiting to be assigned to a compression set.

BFR: “Galaxies” Picture




Discard set (DS): Close enough to a centroid to be summarized
Compression set (CS): Summarized, but not assigned to a cluster
Retained set (RS): Isolated points

Summarizing Sets of Points

For each cluster, the discard set (DS) is summarized by:

- ◆ The number of points, ***N***
- ◆ The vector ***SUM*** of length **d dimensions**, whose i^{th} component is **the sum of the coordinates of the points in the i^{th} dimension**
- ◆ The vector ***SUMSQ*** of length **d dimensions**, whose i^{th} component is the sum of squares of coordinates in i^{th} dimension.



A cluster.


All its points are in the **DS**.

The centroid²⁶

Summarizing Points: Comments

- ◆ $2d + 1$ values represent any size cluster
 - Two vectors of length d = number of dimensions
 - One integer N
- ◆ Average in **each dimension** (**the centroid**) can be calculated as SUM_i / N
 - $\text{SUM}_i = i^{\text{th}}$ component of SUM
- ◆ Variance of a cluster's discard set in dimension i is:
 $(\text{SUMSQ}_i / N) - (\text{SUM}_i / N)^2$
 - And standard deviation is the square root of that (σ_i)
- ◆ **Next step: Actual clustering.**

Note: Dropping the “axis-aligned” clusters assumption would require storing full covariance matrix to summarize the cluster. So, instead of SUMSQ being a d -dim vector, it would be a $d \times d$ matrix, which is too big!



BFR Algorithm

Processing a “Memory-Load” of points:

- ◆ 0) From the initial load we selected the initial k centroids
- ◆ 1) Find those points that are “sufficiently close” to a cluster centroid and **add those points to that cluster and the DS**
 - These points are so close to the centroid that they can be **summarized** and then **discarded**
 - **Adjust statistics of the clusters** to account for the new points
 - Add to N_s , SUM_s , $SUMSQ_s$
 - More on “sufficiently close” later...

Discard set (DS): Close enough to a centroid to be summarized.

Compression set (CS): Summarized, but not assigned to a cluster

Retained set (RS): Isolated points.

BFR Algorithm (2)

- ◆ 2) Use any main-memory clustering algorithm to **cluster the remaining points and the old Retained Set (RS)**
 - “Miniclusters” of more than one point are summarized, added to the CS
 - Summarized in same way (SUM, SUMSQ, N)
 - Singleton clusters (outlying points) become the RS
- ◆ 3) Consider **merging miniclusters in the Compression Set (CS)**
 - Miniclusters from the last step
 - Can't merge with a cluster, but **may merge with one another**
 - Use earlier criteria (merge clusters with smallest distance between centroids, or whose resulting cluster has smallest diameter or radius, etc.)
 - More on this later...

Discard set (DS): Close enough to a centroid to be summarized.

Compression set (CS): Summarized, but not assigned to a cluster

Retained set (RS): Isolated points

BFR Algorithm (3)


- ◆ 4) Points that are assigned to a cluster or minicenter are written out with their assignment to secondary memory
 - (discard DS and CS from memory)
- ◆ (Repeat steps 1 to 4 for all memory-sized chunks of input data)
- ◆ 5) If this is the last chunk of input data, decide what to do with remaining CS and RS
 - Can treat them as outliers and not cluster them at all Or
 - Or Merge sets in CS and points in RS with their nearest cluster:
 1. Assign each point in RS to cluster with nearest centroid
 2. Combine each minicenter in the CS with the cluster whose centroid is closest to the centroid of the minicenter.

Discard set (DS): Close enough to a centroid to be summarized.

Compression set (CS): Summarized, but not assigned to a cluster

Retained set (RS): Isolated points

BFR: “Galaxies” Picture



Discard set (DS): Close enough to a centroid to be summarized.

Compression set (CS): Summarized, but not assigned to a cluster


Retained set (RS): Isolated points

A Few Details...

- ◆ Q1) How do we decide if a point is “close enough” to a cluster that we will add the point to that cluster?
 - New distance measure
- ◆ Q2) How do we decide whether two compressed sets (CS) deserve to be combined into one?

How Close is Close Enough?

- ◆ Q1) We need a way to decide whether to put a new point into a cluster (and discard)
- ◆ The **Mahalanobis distance** is less than a threshold
- ◆ **High likelihood of the point belonging to cluster with the currently nearest centroid.**



Mahalanobis Distance

- ◆ Normalized Euclidean distance from centroid
- ◆ Cluster C has centroid $(\mathbf{c}_1, \dots, \mathbf{c}_d)$ and standard deviations $(\sigma_1, \dots, \sigma_d)$
- ◆ Point $P = (p_1, \dots, p_d)$
 1. Normalize in each dimension: $(p_i - c_i) / \sigma_i$
where σ_i = standard deviation of points in the cluster in the i^{th} dimension
 2. Take sum of squares
 3. Take square root.

$$\sqrt{\sum_{i=1}^d \left(\frac{p_i - c_i}{\sigma_i} \right)^2}$$

Example 7.9

Example 7.9

Suppose a cluster consists of the points (5,1), (6,-2), (7,0)

- Average in **each dimension** (**the centroid**)
can be calculated as **SUM_i / N**
- Variance of a cluster's discard set in dimension i is:
 $(\text{SUMSQ}_i / N) - (\text{SUM}_i / N)^2$
- Standard variation?

Example 7.9 **Standard variation?**

Suppose a cluster consists of the points (5,1), (6,-2), (7,0)

Then $N=3$

$SUM = [18, -1]$

$SUMSQ = [110, 5]$

- Average in **each dimension (the centroid)** can be calculated as SUM_i / N
- Variance of a cluster's discard set in dimension i is:
 $(SUMSQ_i / N) - (SUM_i / N)^2$

The centroid is $SUM/N = [6, -1/3]$

In the first dimension, the variance is

$$110/3 - (18/3)^2 = 0.667$$

The standard variation = $\sqrt{0.667} = 0.816$

In second dimension, the variance is

$$5/3 - (-1/3)^2 = 1.56$$

So the standard variation is 1.25.


Mahalanobis Distance

To decide whether to assign point p to a cluster:

- ◆ **Compute Mahalanobis distance between p and each of the cluster centroids**
- ◆ Choose cluster whose centroid has smallest Mahalanobis distance from p
- ◆ **If this distance is less than a threshold, add p to the cluster**
 - E.g., threshold might be 2 or 4 standard deviations from mean
 - If points are truly normally distributed, very small probability we will not include point that should be in the cluster.

Mahalanobis Distance


- ◆ If clusters are normally distributed in d dimensions, then after transformation, one standard deviation = \sqrt{d}
 - i.e., 68% of the points of the cluster will have a Mahalanobis distance $< \sqrt{d}$
- ◆ Accept a point for a cluster if its M.D. is $<$ some threshold, e.g. **2** standard deviations



Mahalanobis Acceptance Criterion

- Suppose point P is one standard dimension away from centroid in each dimension
 - Each $y_i = 1$ and so the MD of P is \sqrt{d}

We have d dimensions



68% of points have $MD \leq \sqrt{d}$

95% of points have $MD \leq 2\sqrt{d}$


99% of points have $MD \leq 3\sqrt{d}$

Accept a point for a cluster if its M.D. is $<$ some threshold,
e.g. **2** standard deviations


Picture: Equal M.D. Regions

◆ Euclidean vs. Mahalanobis distance


Contours of equidistant points from the origin



Uniformly distributed points,
Euclidean distance



Normally distributed points,
Euclidean distance




Normally distributed points,
Mahalanobis distance

Should 2 CS clusters be combined?

Q2) Should 2 CS miniclusters be combined?

- ◆ Another method (in addition to earlier ones)
- ◆ Compute the variance of the combined minicluster
 - ***N***, ***SUM***, and ***SUMSQ*** allow us to make that calculation quickly
- ◆ Combine if the combined variance is below some threshold
- ◆ **Many alternatives:** Treat dimensions differently, consider density.




The CURE Algorithm


**Extension of k -means to clusters
of arbitrary shapes**

Problem with BFR/*k*-means


- Assumes clusters are normally distributed in each dimension
- And axes are fixed – ellipses at an angle are *not OK*



OK




Not OK




Not OK

The CURE Algorithm




◆ CURE (Clustering Using REpresentatives):

- Assumes a Euclidean space
- Allows clusters to assume any shape
- **Uses a collection of representative points to represent clusters.**



Example: University Salaries




Starting CURE


2 Pass algorithm. Pass 1:

- ◆ **0) Pick a random sample of points that fit in main memory**
- ◆ **1) Initial clusters:**
 - Cluster these points hierarchically – group nearest points/clusters
- ◆ **2) Pick representative points:**
 - For each cluster, pick a sample of points, as dispersed as possible
 - From the sample, pick representatives by moving them (say) 20% toward the centroid of the cluster.


Example: Initial Clusters



Example: Pick Dispersed Points




Example: Pick Dispersed Points



Finishing CURE

Pass 2:

- ◆ Now, rescan the whole dataset and visit each point p in the data set
- ◆ Place it in the “closest cluster”
 - Normal definition of “closest”:
Find the closest representative to p and assign it to representative’s cluster.



Example: A Circle and a Ring Cluster




Figure 7.12: Two clusters, one surrounding the other

Starting CURE

2 Pass algorithm

Pass 1:

- ◆ 0) Pick random sample of points that fits in main memory
- ◆ 1) Initial clusters:
 - Cluster these points **hierarchically** – group nearest points/clusters
 - **Note: don't use distance between centroids: centroids are same!**
- ◆ 2) Pick representative points:
 - For each cluster, pick a sample of **representative points**, as dispersed as possible
- ◆ 3) Move representative points:
 - From the sample, **move representatives by (say) 20% toward the centroid of the cluster**
 - Note: requires a Euclidean space.

Pick Representative Points in Clusters




Figure 7.13: Select representative points from each cluster, as far from one another as possible

Move Representative Points 20% of the Distance to Centroid




Figure 7.14: Moving the representative points 20% of the distance to the cluster's centroid

Note: both clusters have the same centroid!

Finishing CURE

Pass 2:

- ◆ Now, **rescan the whole dataset** and visit each point p in the data set
- ◆ **Place each point in the “closest cluster”**
 - Normal definition of “**closest**”: Find the **closest representative to p** and **assign it to the representative’s cluster**
- ◆ **End!**

Notes on this Example

- ◆ Could argue that the ring and the circle should be merged into **one cluster**, since their centroids are the same
- ◆ Whether we have one cluster or two depends on:
 - **Fraction of the distance to the centroid that we move the representative points**
 - **Choice of how far apart representative points of two clusters need to be to avoid merger**

MapReduce for Clustering Algorithms

- ◆ Data Clustering using *MapReduce*: A Look at Various *Clustering Algorithms* Implemented with *MapReduce* Paradigm
- ◆ By Varad Meru
- ◆ <http://www.slideshare.net/VaradMeru/data-clustering-using-map-reduce>

MapReduce for Clustering Algorithms

- ◆ Start (as usual) with k points representing initial cluster centroids
- ◆ **Mappers:**
 - Receive portion of the data points and k centroids
 - **Do distance calculations** of each point from centroids
 - **Assign each point to its closest centroid**
 - **Emit (clusterID, datapoint)**
- ◆ **Reducers:**
 - Work on a clusterID, list of data points in the cluster
 - **Computes new centroid** for the cluster
 - **Writes centroid to a new centroid file** for next MapReduce phase
- ◆ **Termination:**
 - After specified **number of iterations** or **when clusters stabilize**.

Summary

- ◆ **Clustering:** Given a **set of points**, with a notion of **distance** between points, **group the points** into some number of *clusters*
- ◆ **Algorithms:**
 - Agglomerative **hierarchical clustering**:
 - Centroid (for Euclidean spaces)
 - Clustroid (for non-Euclidean spaces)
 - **k-means**:
 - Point assignment; Euclidean space; initialization, picking k
 - **BFR**
 - Extends k-means for data sets that don't fit in memory; assumes clusters are normally distributed about a centroid
 - **CURE**
 - Does not restrict the shape of clusters; represent clusters by a collection of representative points.