

# Web Advertising

Thanks for source slides and material to: J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmms.org>

# Chapter 8 Overview

- Ability of all sorts of Web applications to support themselves through advertising
- **Most lucrative venue for on-line advertising: SEARCH**
- **Adwords model (Google): matching search queries to advertisements**
  - Algorithms for optimizing this assignment
  - **Greedy** algorithms
  - **Online** algorithms
- Selecting items to advertise at an on-line store
  - Use collaborative filtering.

# Types of Web Ads

- **Advertisers post ads directly**
  - Craig's List, auto trading sites
- **Advertisers pay for display ads to be placed on Web sites**
  - **Fixed price per *impression*** (one display of the ad with download of page by a user)
- **Online stores show ads**
  - Amazon, Macy's, etc.
  - Not paid for by manufacturers of product advertised
  - **Selected by store to maximize probability customer will buy product**
  - **Collaborative Filtering**
- **Search ads are placed among results of a search query**
  - Advertisers bid for right to have their ad shown in response to certain queries
  - **Pay only if ad is clicked on.**

# Online Algorithms

- **Classic model of algorithms**
  - You get to see the entire input, then compute some function of it
  - In this context, “offline algorithm”,
- **Online Algorithms**
  - You get to see the input one piece at a time, and need to make irrevocable decisions along the way
  - **Make decisions without knowing the future**
  - **For search: only know past queries and current query; don't know what queries will come in later**
  - Similar to handling data streams,
- **An online algorithm cannot always do as well as an offline algorithm.**

# Example 8.1

- **Knowing the future could help**
- **Manufacturer A** of antique furniture **bids 10 cents** on search term **“chesterfield”**
- **Manufacturer B** of conventional furniture **bids 20 cents** on both terms **“sofa” and “chesterfield”**
- Both have monthly **budget of \$100**
  - **A can place its ad 1000 times, B can place its ad 500 times**
- Query “chesterfield” arrives
- Can only display one ad
- Might display B’s ad because B bid more
- However, if there are many queries for “sofa” and few for “chesterfield,” A will never spend its full budget
- Sending “chesterfield” queries to A might increase overall revenue
- **Without knowing the future, on-line algorithm may not perform as well as offline.**

# Greedy Algorithms

- **Make their decision in response to each input element by maximizing some function of (input element, past)**
- Example 8.2: Greedy algorithm would **assign the query to the highest bidder who still has budget left**
- ◆ Recall: **A** bid 10 cents on “chesterfield”, **B** bid 20 cents on “sofa”, chesterfield”

## Worst case:

- 500 “chesterfield” queries arrive
- **Greedy algorithm: all are assigned to B with the higher bid**
- **B budget used up; earns \$100 revenue**
- **Followed by 500 “sofa” queries that will get no ad assigned: no revenue**

## Optimal:

- **First 500 chesterfield queries assigned to A, earn \$50**
- **Next 500 sofa queries assigned to B, earn \$100**
- **Total revenue \$150.**

# Online Bipartite Matching

# The Matching Problem

- Simplified version of the problem of matching ads to search queries
- “Maximal matching”: involves bipartite graphs with two sets of nodes
- All edges connect node on left set to node in right set.





Figure 8.1: A bipartite graph




# Example: Bipartite Matching



**Nodes: Queries and Ads**


**Goal: Match queries to ads so that maximum number of matchings are made**

# Example: Bipartite Matching



$M = \{(1,a), (2,b), (3,d)\}$  is a **matching**  
Cardinality of matching =  $|M| = 3$

# Example: Bipartite Matching



$M = \{(1,c), (2,b), (3,d), (4,a)\}$  is a  
**perfect matching**

**Perfect matching:** all vertices of the graph are matched  
**Maximal matching:** a matching that contains the largest possible number of matches


# Matching Algorithm

- **Problem:** Find a maximal matching for a given bipartite graph
  - A perfect one if it exists
- There is a polynomial-time offline algorithm based on augmenting paths (Hopcroft & Karp 1973, see [http://en.wikipedia.org/wiki/Hopcroft-Karp\\_algorithm](http://en.wikipedia.org/wiki/Hopcroft-Karp_algorithm))
- **But what if we do not know the entire graph upfront?**

# Online Graph Matching Problem

- Initially, we are given the set ads
- In each round, one set of query terms is added
  - Relevant edges are revealed
  - Indicate which advertisers have bid on those query terms
- At that time, we have to decide to either:
  - Pair the query with an ad
  - Do not pair the query with any ad.

# Online Graph Matching: Example



(1,a)  
(2,b)  
(3,d).

# Greedy Algorithm

## ➤ Greedy algorithm for the online graph matching problem:

- Pair the new query with **any** eligible ad
  - If there is none, do not pair query

## ➤ How good is the algorithm?

# Competitive Ratio

- For input  $I$ , suppose greedy produces matching  $M_{greedy}$  while an optimal matching is  $M_{opt}$

Competitive ratio =

$$\min_{\text{all possible inputs } I} (|M_{greedy}| / |M_{opt}|)$$


(what is greedy's worst performance over all possible inputs  $I$ ).



# Analyzing the Greedy Algorithm

- Consider a case:  $M_{greedy} \neq M_{opt}$
- Consider the set  $Q$  of queries matched in  $M_{opt}$  but not in  $M_{greedy}$
- $A$  is the set of ads that are adjacent to a non-matched query in  $Q$  that are already matched in  $M_{greedy}$ 
  - Why: If there exists such a non-matched (by  $M_{greedy}$ ) ad adjacent to a non-matched query, then greedy would have matched them
- Since ads  $A$  are already matched in  $M_{greedy}$  then
 

(1)  $|M_{greedy}| \geq |A|$



# Analyzing the Greedy Algorithm

## ➤ Summary so far:

- Queries  $Q$  matched in  $M_{opt}$  but not in  $M_{greedy}$
- (1)  $|M_{greedy}| \geq |A|$

## ➤ There are at least $|Q|$ such ads in $A$ ( $|Q| \leq |A|$ ) otherwise the optimal algorithm couldn't have matched all queries in $Q$

- So: (2)  $|Q| \leq |A| \leq |M_{greedy}|$

## ➤ By definition of $Q$ also: (3) $|M_{opt}| \leq |M_{greedy}| + |Q|$


- Worst case is when  $|Q| = |A| = |M_{greedy}|$

## ➤ Combining (2) and (3)


## ➤ $|M_{opt}| \leq 2|M_{greedy}|$ then $|M_{greedy}|/|M_{opt}| \geq \frac{1}{2}$

## ➤ Competitive Ratio = $\frac{1}{2}$

## ➤ Greedy's worst performance over all possible inputs $\frac{1}{2}$ .



# Worst-case Scenario



- **Worst case** is when  $|Q| = |A| = |M_{greedy}|$
- $Q = \{c, d\}$  – queries with no matching ad
- $A = \{1, 2\}$  – ads that are adjacent to a query in  $Q$  but are already matched to another query
- $|M_{greedy}| = 2, |Q| = 2, |A| = 2$
- **Optimal matching: (1,c), (2,d), (3,b), (4,a)**
- $|M_{opt}| = 4$
- $|M_{greedy}| / |M_{opt}| = \frac{1}{2}$  (competitive ratio)

# Summary

- Web Advertising
- Matching search queries to advertisements
- Matching Algorithm
  - Algorithms for optimizing this assignment
  - Greedy algorithms
  - Online algorithms
- competitive ratio
  - Competitive ratio for Greedy algorithms =  $1/2$